

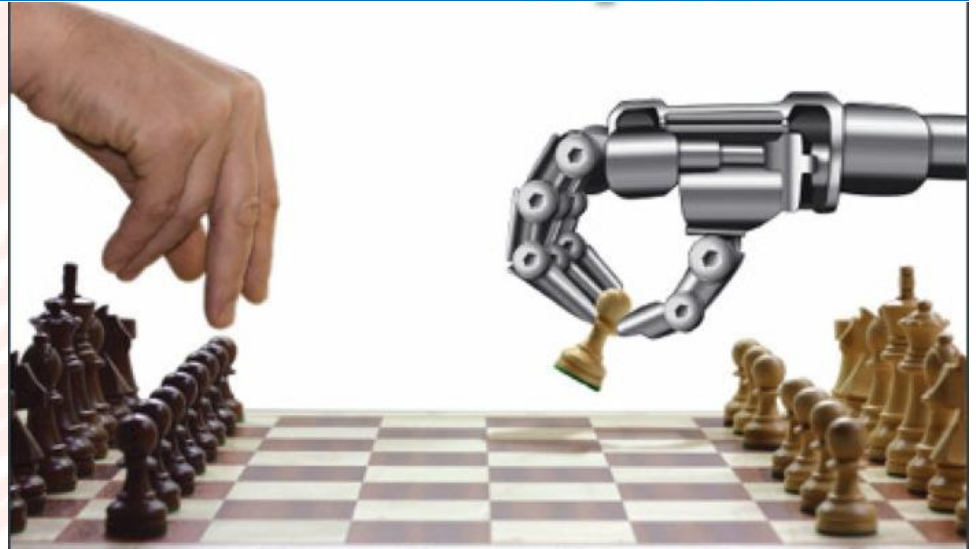


FACULTY OF INFORMATION TECHNOLOGY

Artificial Intelligence Fundamentals (NM TTNT)

Semester 1, 2023/2024

Chapter 5. Adversarial Search



Content

- ▶ What are games?
- ▶ Optimal decisions in games
 - Which strategy leads to success?
- ▶ α - β pruning
- ▶ Games of imperfect information
- ▶ Games that include an element of chance

What are and why study games?

- ▶ Games are a form of multi-agent environment
 - What do other agents do and how do they affect our success?
 - Cooperative vs. competitive multi-agent environments.
 - Competitive multi-agent environments give rise to adversarial problems a.k.a. games
- ▶ Game playing is a good problem for AI research

What are and why study games?

- ▶ Game playing is non-trivial
 - Players **need "human-like" intelligence**
 - Games can be very **complex** (e.g. chess, go)
 - Requires **decision making within limited time**
- ▶ Games usually are:
 - Well-defined and repeatable
 - Limited and accessible

Types of Games

► Perfect Information (fully observable)

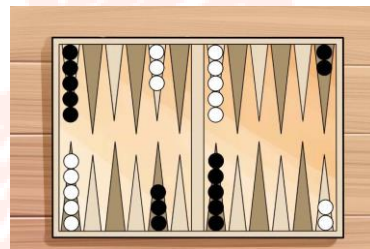
◦ Deterministic

- Chess,
- Checkers,
- Go,
- Othello



◦ Chance

- Backgammon,
- Monopoly



Types of Games

► Imperfect Information (partially observable)


- Deterministic

- Stratego
- Battleship



- Chance

- Brigde,
- Poker,
- Scrabble,
- Nuclear War



Types of Games

- ▶ This course focuses on:
 - Perfect Information (fully observable)
 - Deterministic



Relation of Games to Search

- ▶ Solution is (heuristic) method for finding goal
- ▶ **Heuristics** and **CSP** (Constraint Satisfaction Problems) techniques can find *optimal* solution
- ▶ **Evaluation function**: estimate of **cost from start to goal** through given node
- ▶ Examples: path planning, scheduling activities
- ▶ Solution is strategy (specifies move for every possible opponent reply).
- ▶ Time limits force an approximate solution
- ▶ **Evaluation function**: evaluate "**goodness**" of game position
- ▶ Examples: chess, checkers, Othello, backgammon

Search – no adversary

Games– adversary

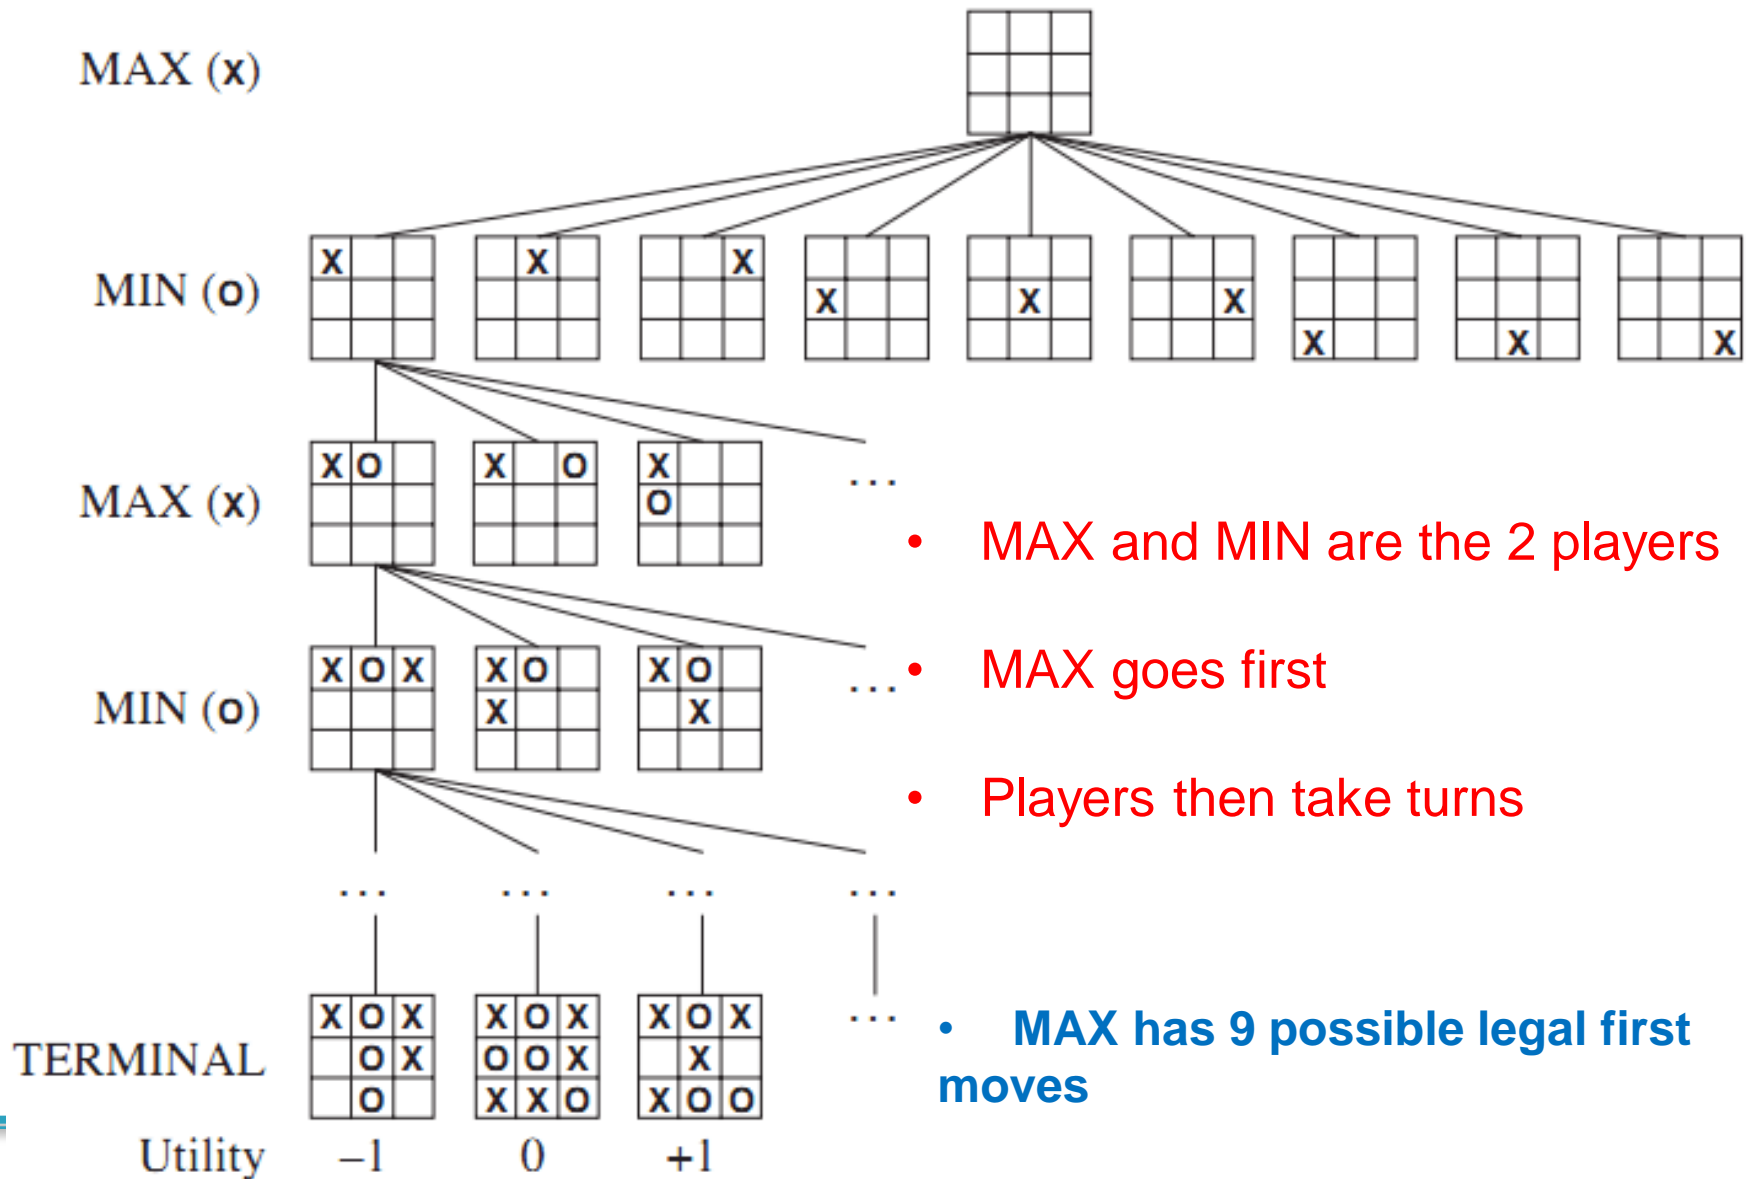
Game setup

- ▶ Two players: *MAX* and *MIN*
- ▶ *MAX* moves first and they take turns until the game is over.
 - Winner gets award,
 - Loser gets penalty.
- ▶ MAX uses search tree to determine next move.

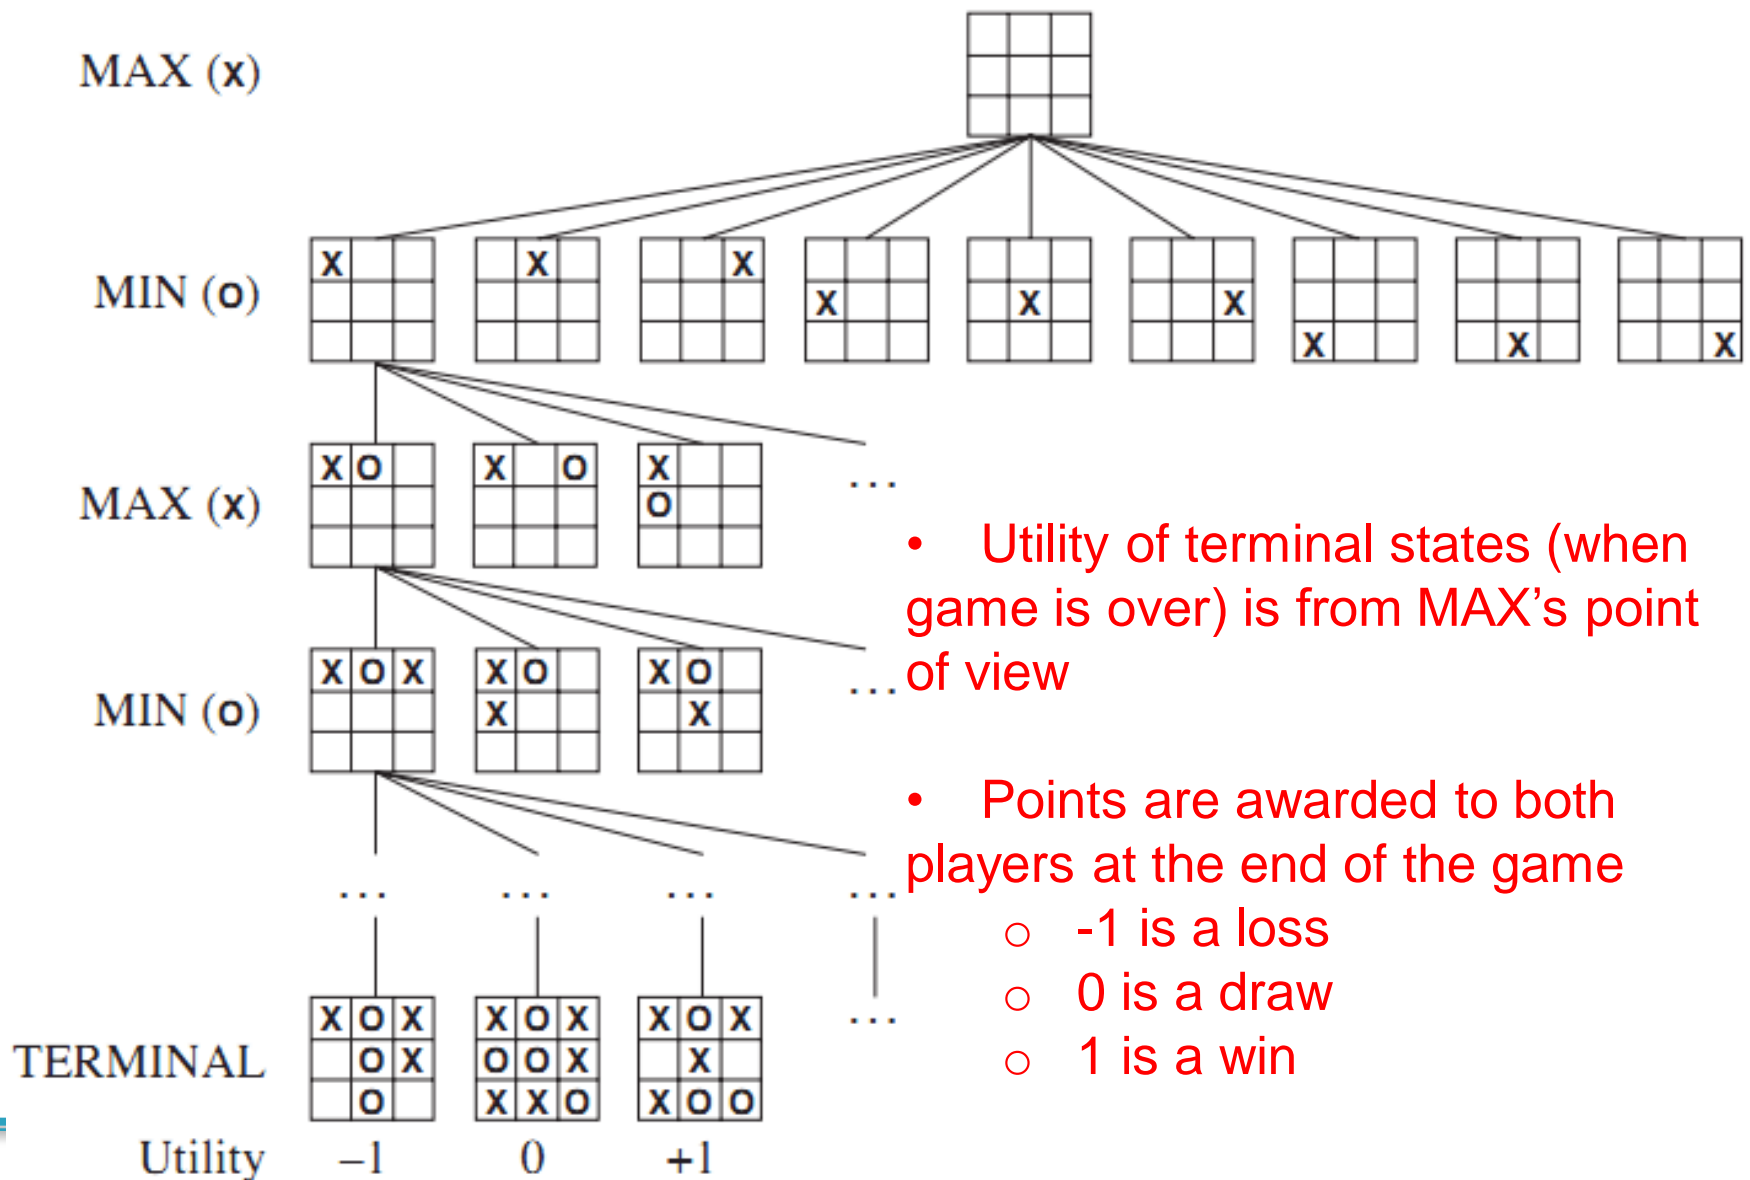
Game Search

- ▶ Problem Formulation:
 - **States**: board configuration of chess
 - **Successor function**: legal moves a player can make.
 - **Goal test**: determines when the game is over.
 - **initial state**: start board configuration
 - **Utility function**: measures the outcome of the game and its desirability
- ▶ Search objective:
 - Find **the sequence of player's decisions** (moves) maximizing its utility
 - Consider **the opponent's moves** and **their utility**

Game Tree



Game Tree



Game Playing as Search: Complexity

- ▶ Assume the opponent's moves *can* be predicted given the computer's moves.
- ▶ How **complex** would search be in this case?
 - Worst case: $O(b^d)$, **b** branching factor, **d** depth
 - **Tic-Tac-Toe**: ~5 legal moves, max of 9 moves
 - $5^9 = 1,953,125$ states
 - **Chess**: ~35 legal moves, ~100 moves per game
 - $35^{100} \sim 10^{154}$ states (but “only” $\sim 10^{40}$ legal states)

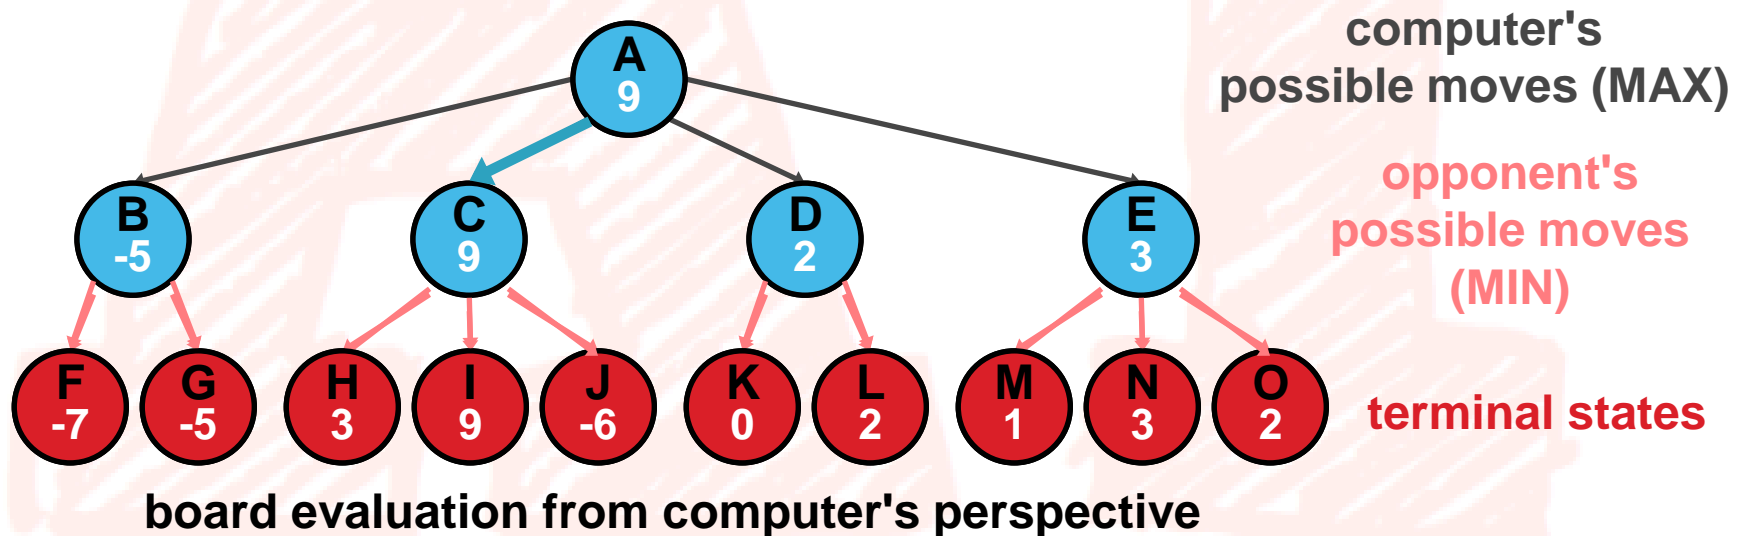
Common games produce enormous search trees!!

Greedy Search Game Playing

- ▶ A *utility* function maps each terminal state of the board to a numeric value corresponding to the value of that state to the computer.
 - **positive for winning**, large positive value: means better for computer (MAX)
 - **negative for losing**, large negative value: means better for opponent (MIN)
 - **zero for a draw**
 - typical values (lost to win):
 - **-infinity to +infinity**
 - **-1.0 to +1.0**

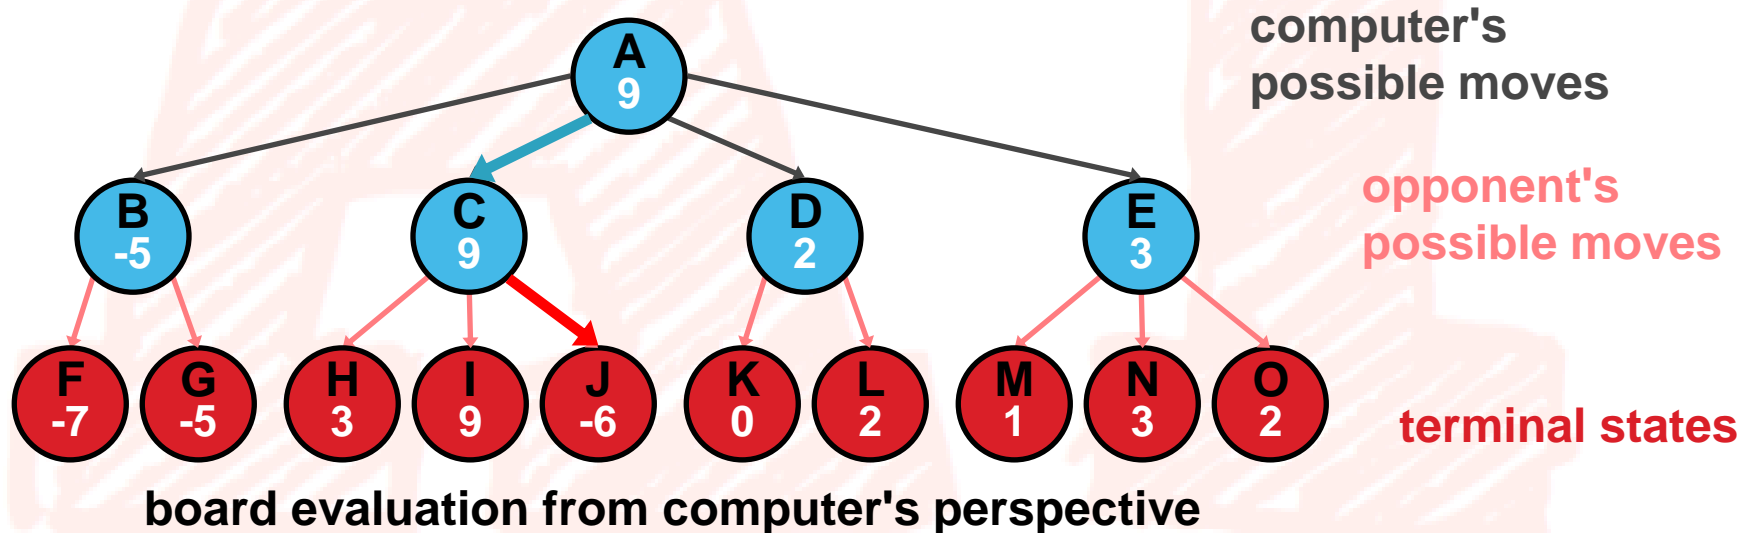
Greedy Search Game Playing

- ▶ Expand each branch to the **terminal states**
- ▶ Evaluate the utility of each terminal state
- ▶ Choose the move that results in the board configuration with **the maximum value**



Greedy Search Game Playing

- ▶ Assuming a reasonable search space, what's the problem with greedy search?
 - It ignores what the opponent might do!
 - e.g. MAX (computer) chooses C.
MIN (opponent) chooses J and defeats computer.



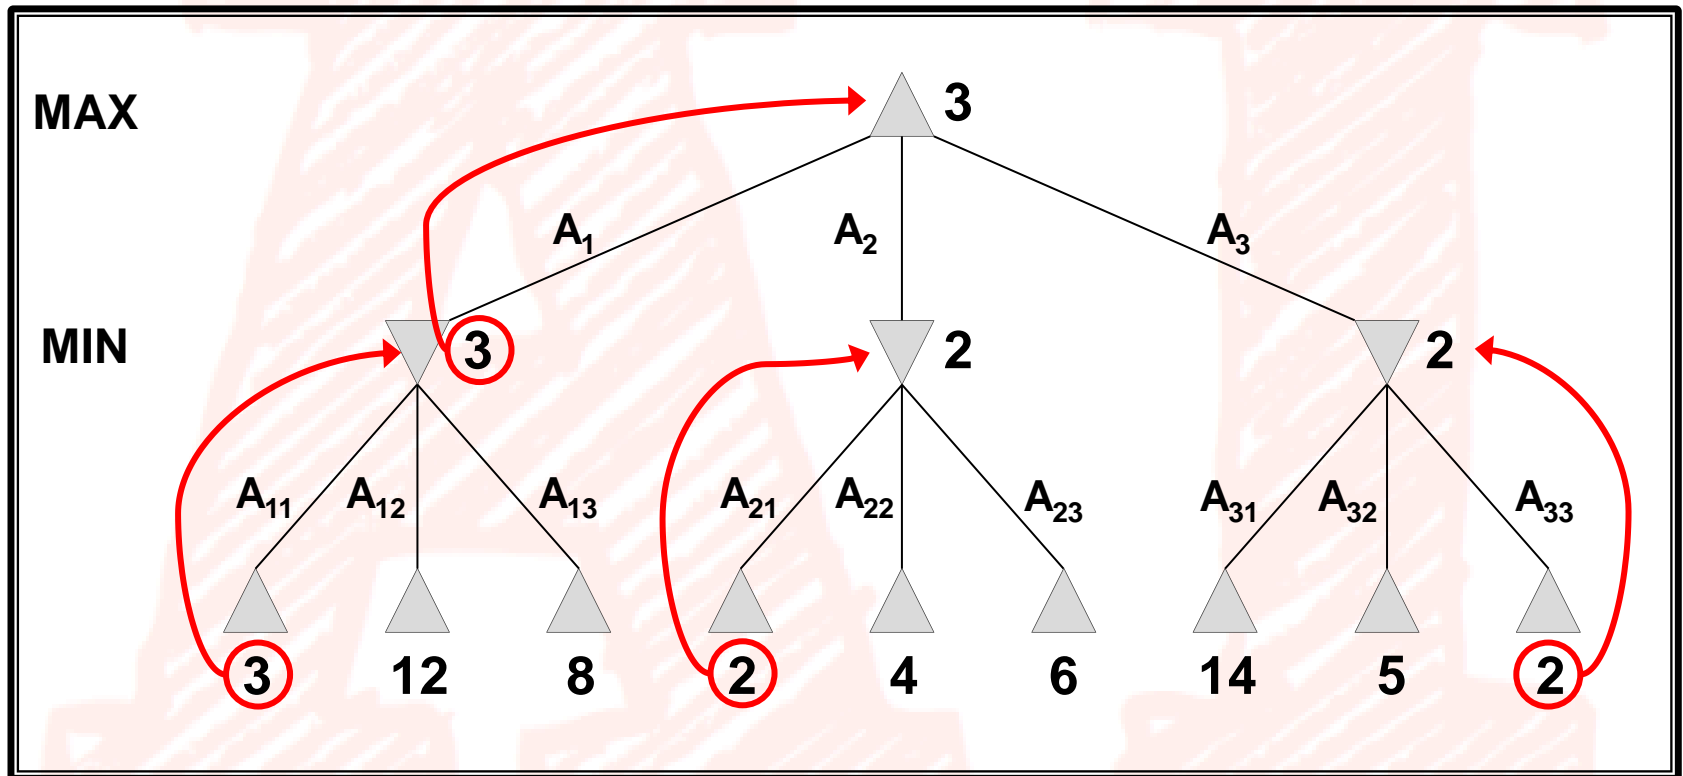
Minimax principle – Optimal strategies

- ▶ Chooses the best move considering both its move and the opponent's best move
- ▶ Assumption: Both players play optimally!!
 - **MAX** (computer) **maximizing** the utility under the assumption after it moves **MIN** (opponent) will choose the **minimizing** move.
- ▶ Given a game tree, the optimal strategy can be determined by using the minimax value of each node:

MINIMAX-VALUE(n) =
UTILITY(n) If n is a terminal
 $\max_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$, If n is a max node
 $\min_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$, If n is a min node

Two-Players Game Tree

The minimax decision



Minimax maximizes the worst-case outcome for max.

What if MIN does not play optimally?

- ▶ Definition of optimal play for **MAX** assumes **MIN** plays optimally: maximizes worst-case outcome for **MAX**.
- ▶ But if **MIN** does not play optimally, **MAX** will do even better. [Can be proved.]

Minimax: Direct Algorithm

For each move by the MAX (computer):

- ▶ Perform **depth-first search to a terminal state**
- ▶ **Evaluate each terminal state**
- ▶ **Propagate upwards the minimax values**
 - if opponent's move minimum value of children backed up
 - if computer's move maximum value of children backed up
- ▶ **choose move with the maximum of minimax values of children**
- ▶ **Note:**
 - minimax values gradually propagate upwards as DFS proceeds: i.e., minimax values propagate up in “left-to-right” fashion
 - minimax values for sub-tree backed up “as we go”, so only **$O(bd)$** nodes need to be kept in memory at any time

Minimax Algorithm

function MINIMAX-DECISION(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

return the *action* in $\text{SUCCESSORS}(\textit{state})$ with value v

function MAX-VALUE(*state*) *returns a utility value*

if $\text{TERMINAL-TEST}(\textit{state})$ **then return** $\text{UTILITY}(\textit{state})$

$v \leftarrow -\infty$

for each s **in** $\text{SUCCESSORS}(\textit{state})$ **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return v

function MIN-VALUE(*state*) *returns a utility value*

if $\text{TERMINAL-TEST}(\textit{state})$ **then return** $\text{UTILITY}(\textit{state})$

$v \leftarrow +\infty$

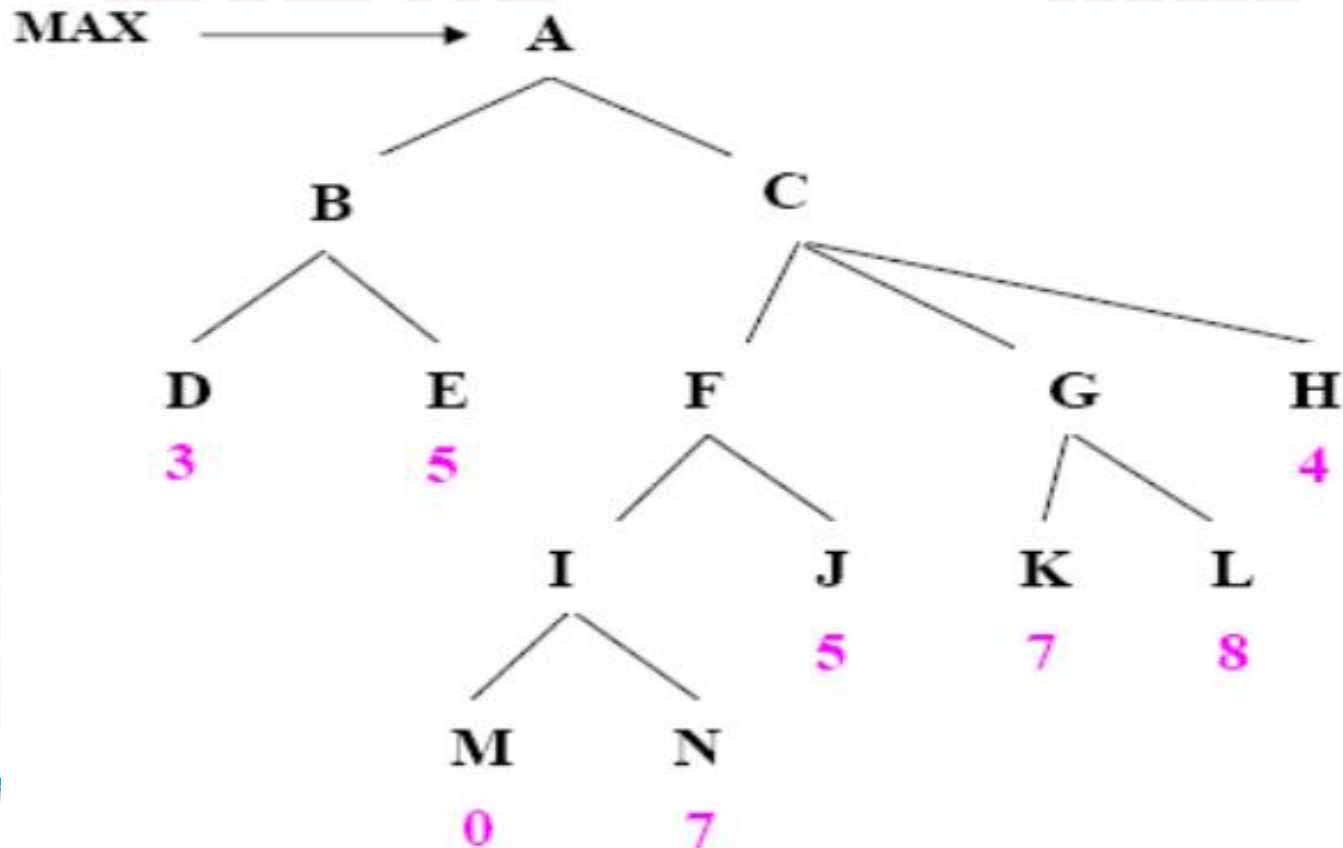
for each s **in** $\text{SUCCESSORS}(\textit{state})$ **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return v

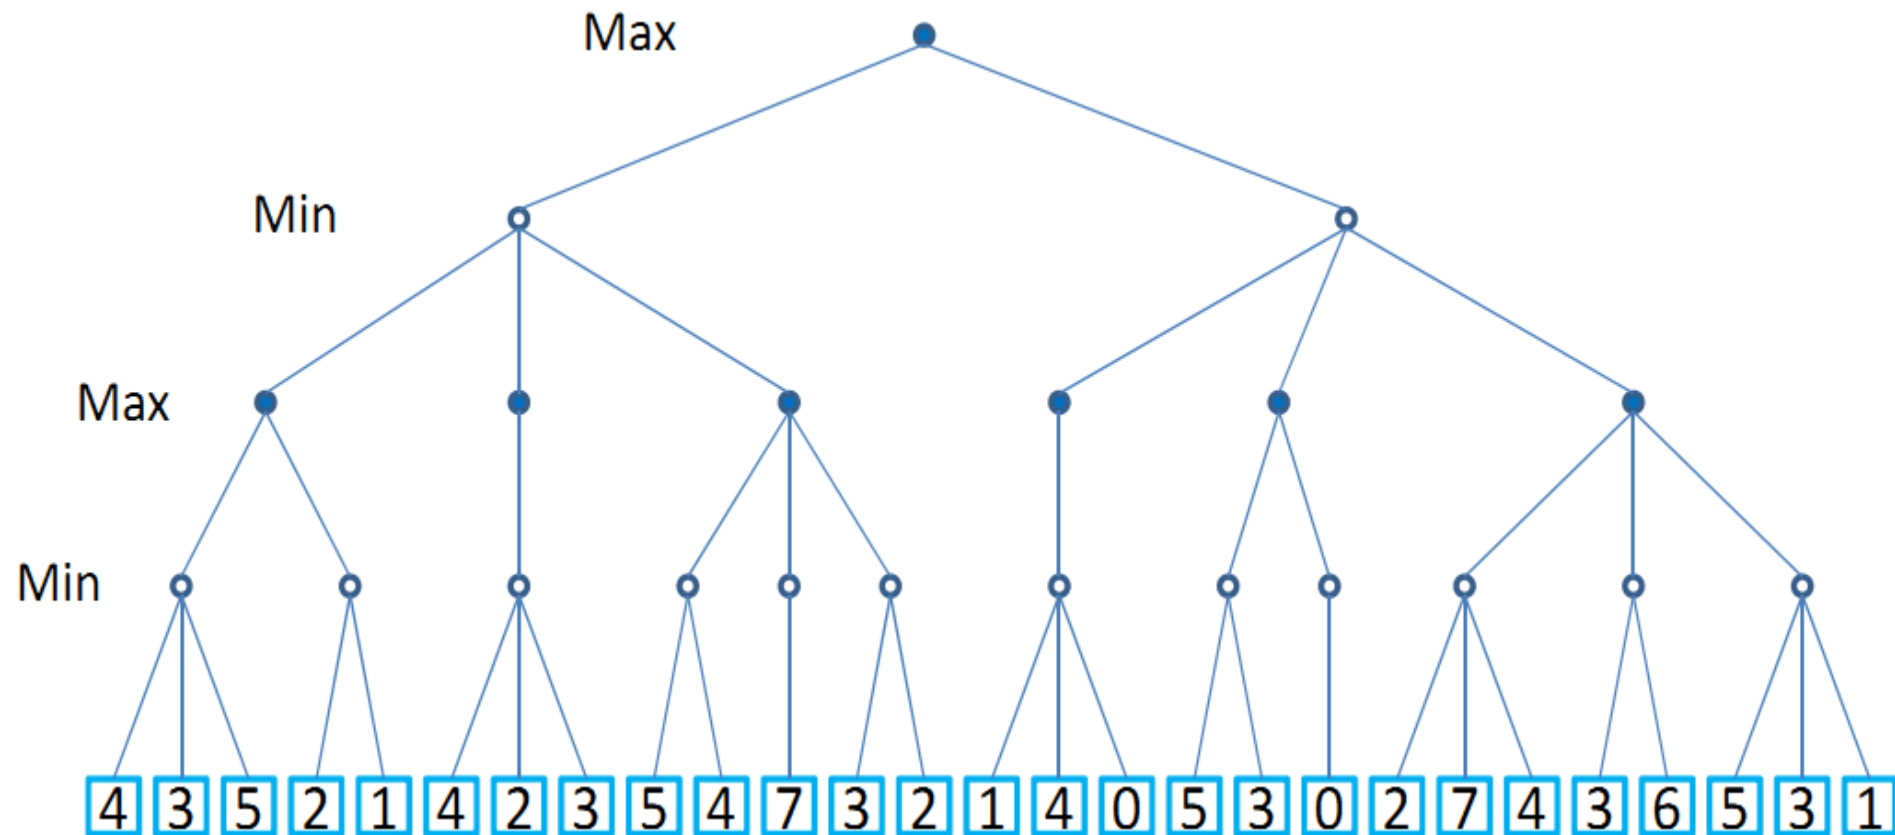
Exercise 1

- ▶ Perform the minimax algorithm on the figure below.



Exercise 2

- Perform the minimax algorithm on the figure below.



Properties of Minimax

Criterion	Minimax
Complete?	Yes (against an optimal opponent)
Time complexity	given branching factor b , $O(b^m)$
Space complexity	$O(bm)$ (depth-first exploration)
Optimal?	Yes (if tree is finite)

- ▶ **Time complexity is a major problem!**
Player typically only has a finite amount of time to make a move!!
- ▶ For chess, $b \approx 35$, $m \approx 100$ for "reasonable" games
→ exact solution completely infeasible

Problem of minimax search

- ▶ Number of games states is **exponential to the number of moves**.
- ▶ Some of the branches of the game tree won't be taken if playing against an intelligent opponent
- ▶ Solution: can "**prune**" those branches from the tree ==> **Alpha-beta pruning**
- ▶ While doing DFS of game tree, keep track of:
 - **Alpha** = Highest value found so far at any choice point along the MAX path
 - Lower bound on node's utility
 - **Beta** = Lowest value found so far at any choice point along the MIN path
 - Higher bound on node's utility

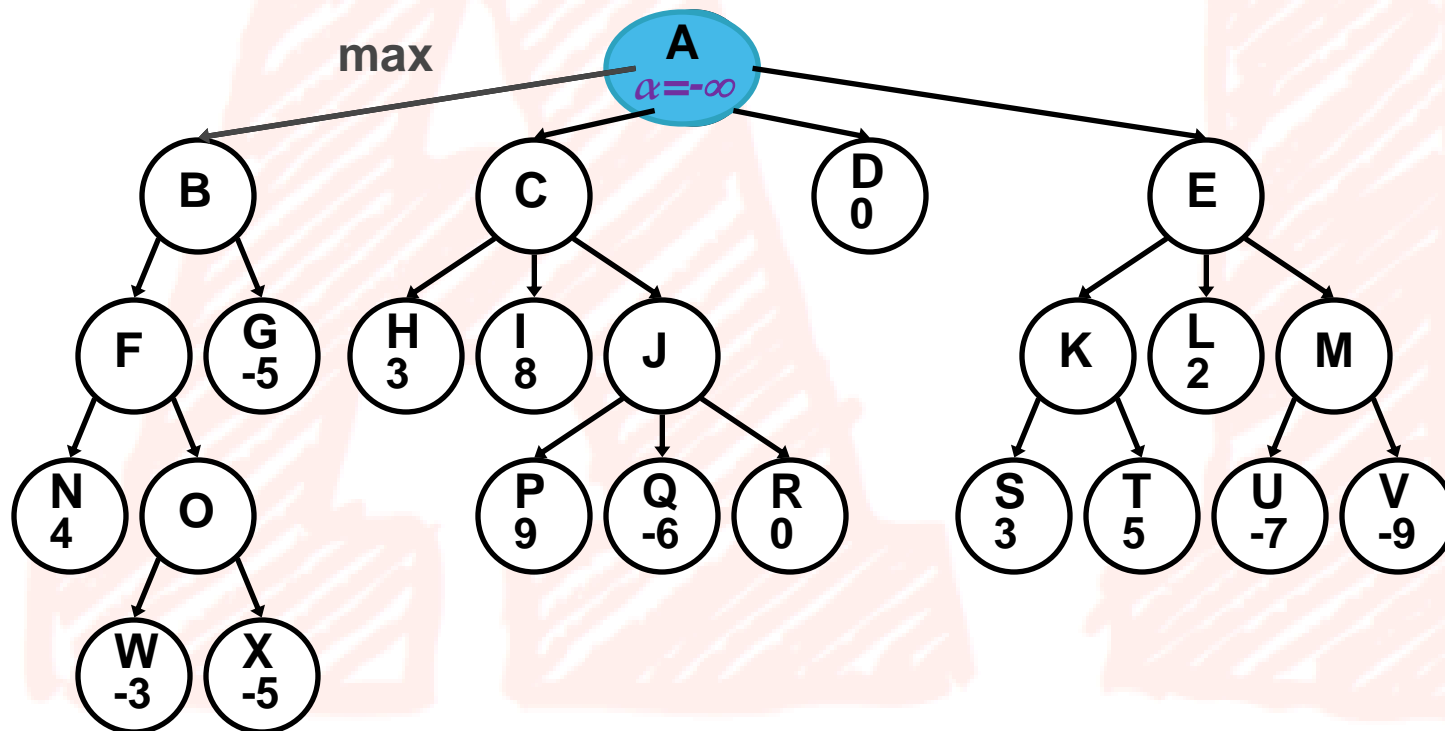
Alpha-Beta pruning

- ▶ Beta cutoff pruning occurs when maximizing (MAX's turn):
 - If $\alpha \geq \text{parent's beta}$, stop expanding
 - Why stop expanding children?
 - Opponent shouldn't allow the MAX to make this move
- ▶ Alpha cutoff pruning occurs when minimizing (MIN's turn):
 - If $\beta \leq \text{parent's alpha}$, stop expanding
 - Why stop expanding children?
 - MAX shouldn't take this route

Alpha-Beta Search Example

$\text{minimax}(A, 0, 4)$ alpha initialized to $-\infty$

Expand A? Yes since there are successors, no cutoff test for root



Call
Stack

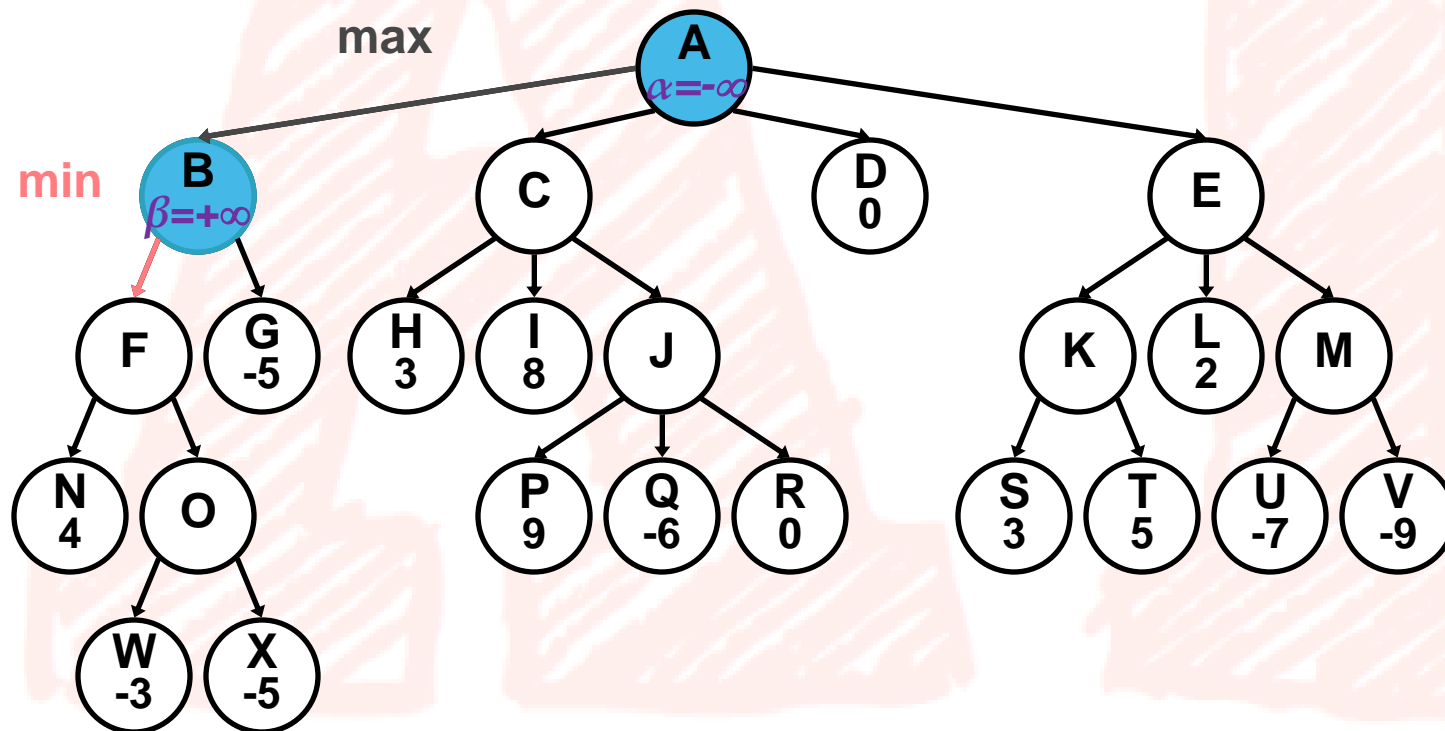
A

Alpha-Beta Search Example

$\text{minimax}(B, 1, 4)$

beta initialized to +infinity

Expand B? Yes since A's alpha \geq B's beta is **false**, no alpha cutoff



Call
Stack

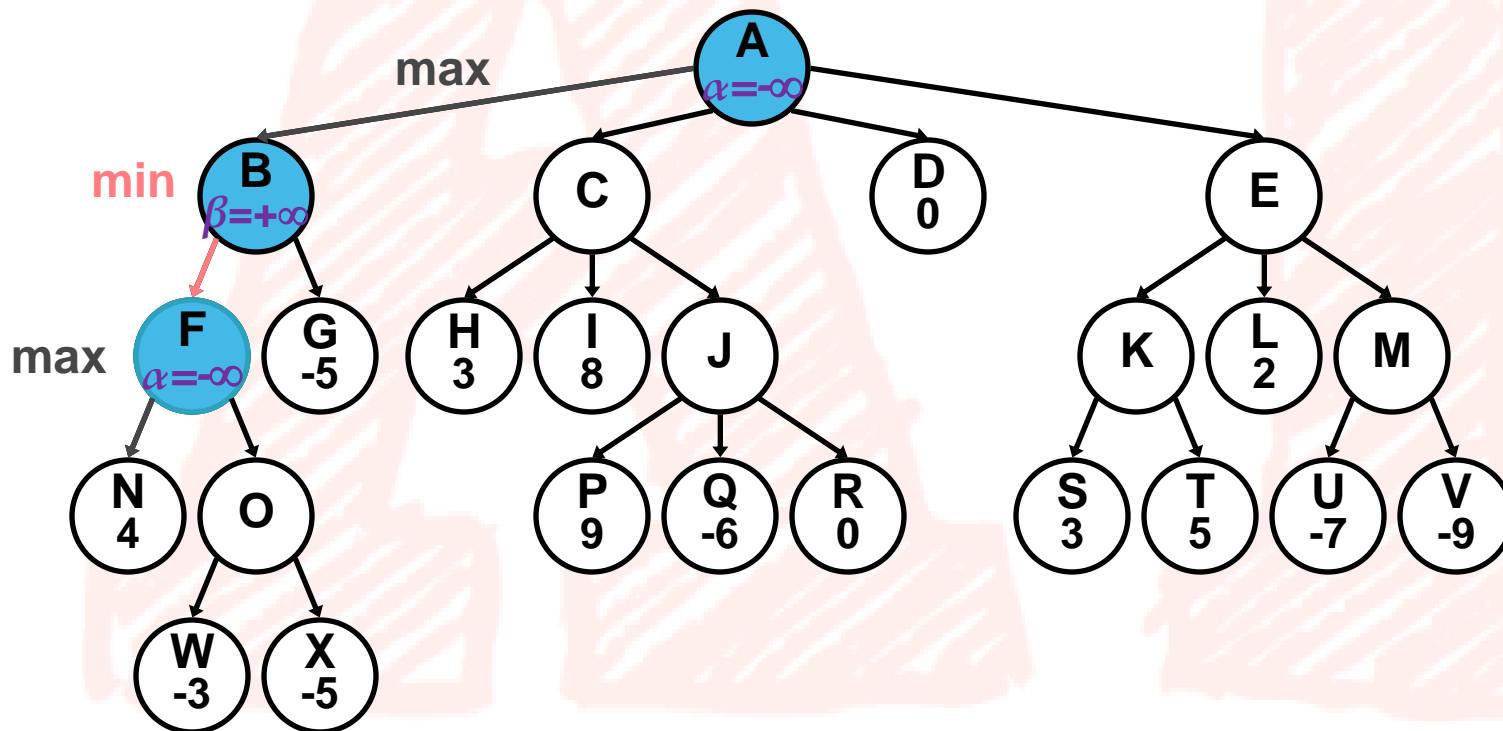
B
A

Alpha-Beta Search Example

$\text{minimax}(F, 2, 4)$

alpha initialized to -infinity

Expand F? Yes since F's alpha \geq B's beta is **false**, no beta cutoff



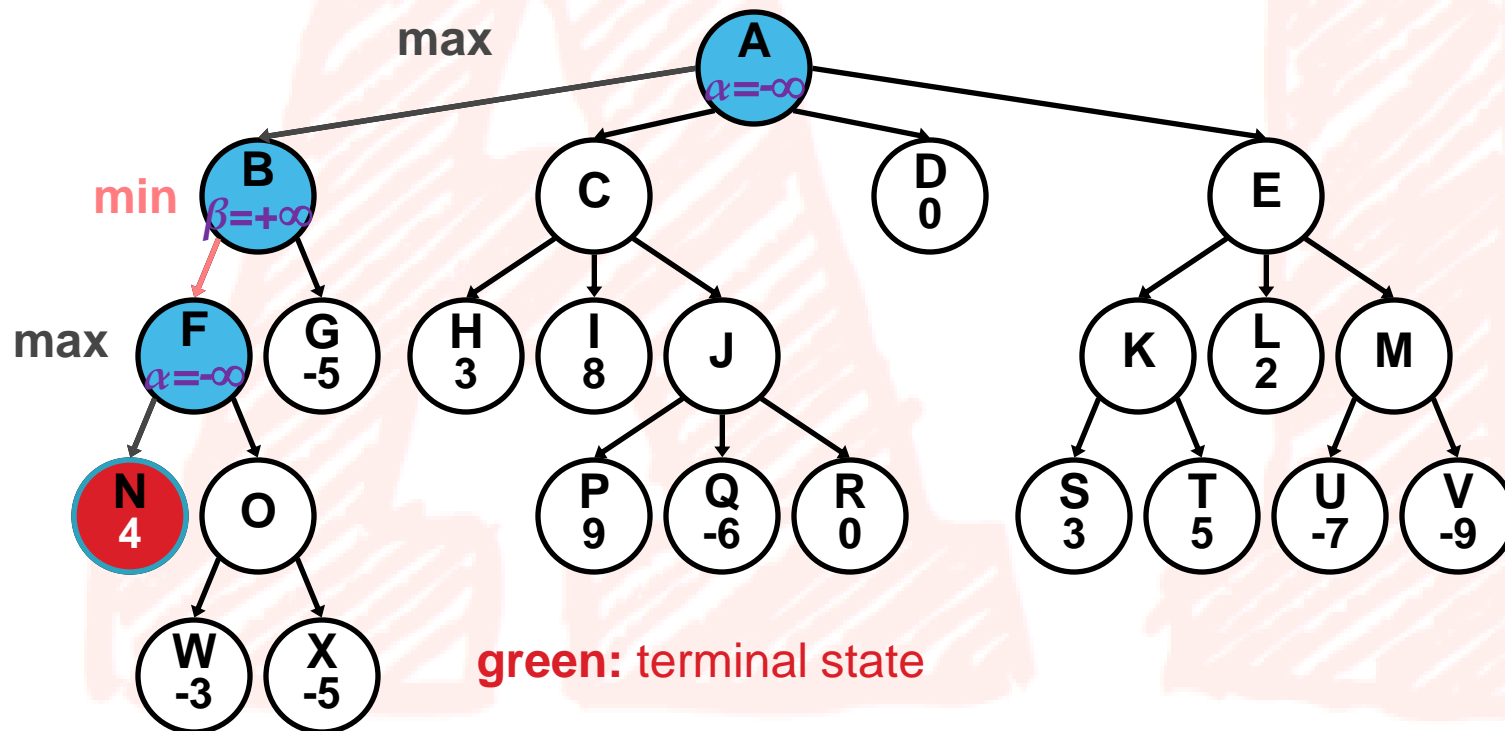
Call
Stack

F
B
A

Alpha-Beta Search Example

$\text{minimax}(\text{N}, 3, 4)$

evaluate and return SBE value



Call
Stack

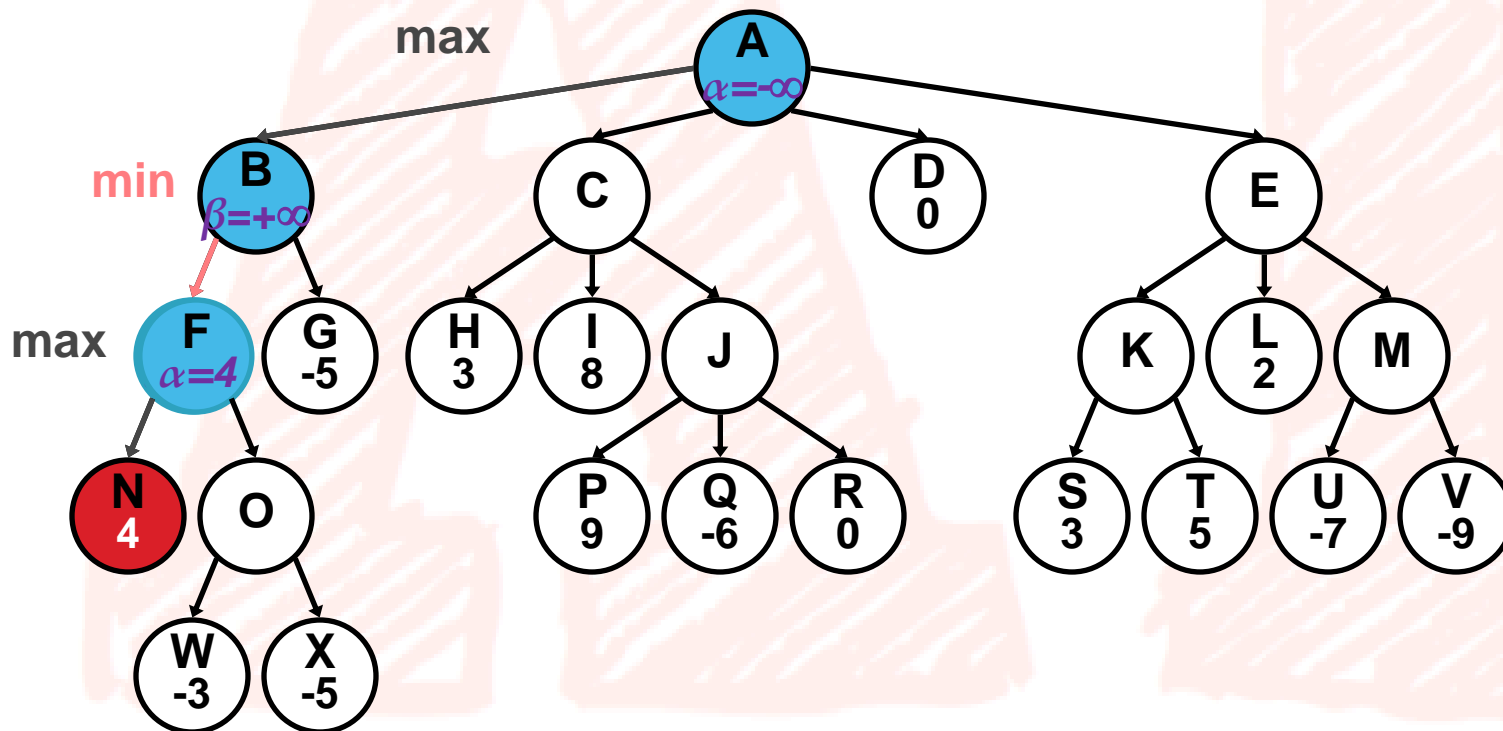
N
F
B
A

Alpha-Beta Search Example

back to
 $\text{minimax}(F, 2, 4)$

$\alpha = 4$, since $4 \geq -\infty$ (maximizing)

Keep expanding F? Yes since F's $\alpha \geq$ B's beta is **false**, no beta cutoff



Call
Stack

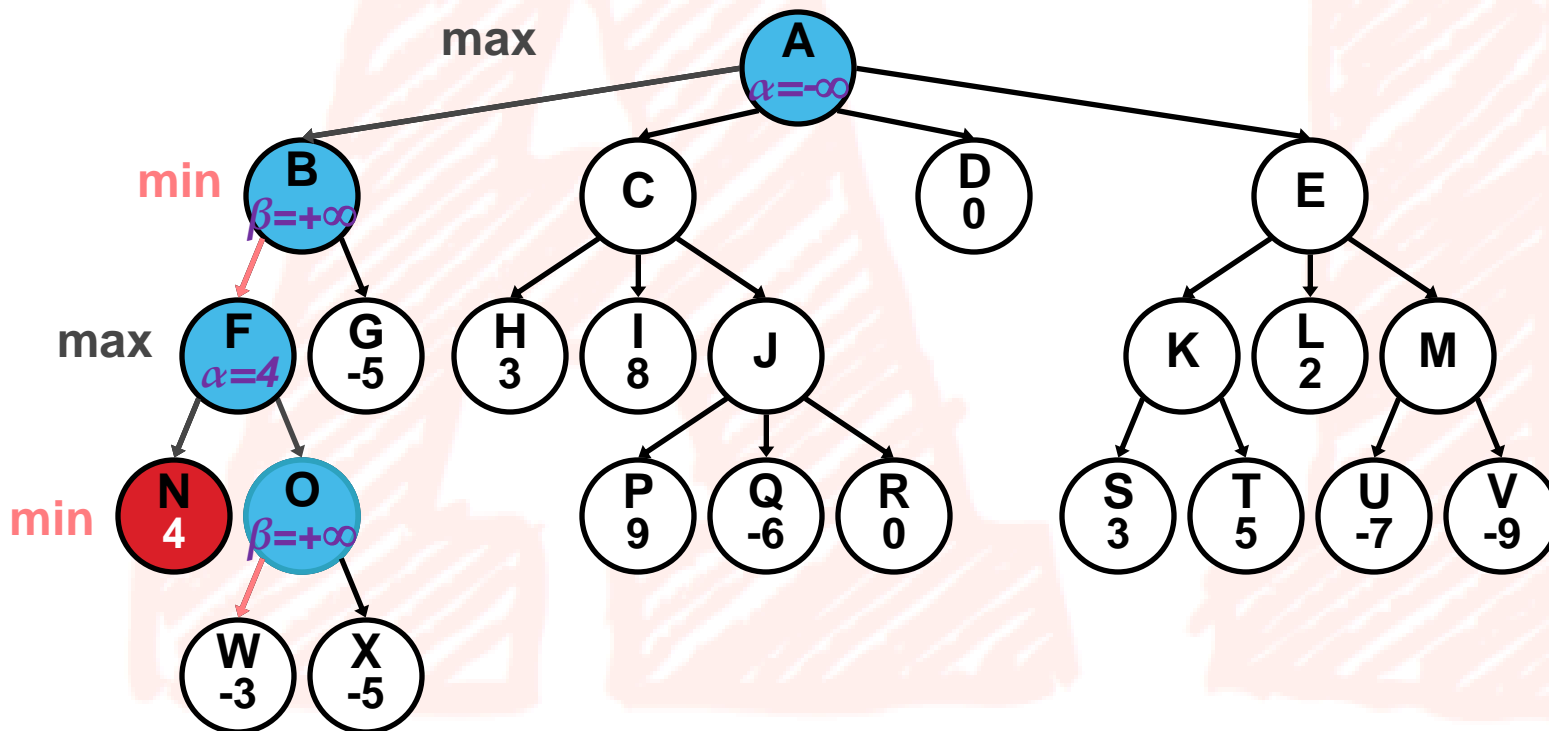
F
B
A

Alpha-Beta Search Example

$\text{minimax}(0, 3, 4)$

beta initialized to +infinity

Expand O? Yes since F's alpha \geq O's beta is **false**, no alpha cutoff



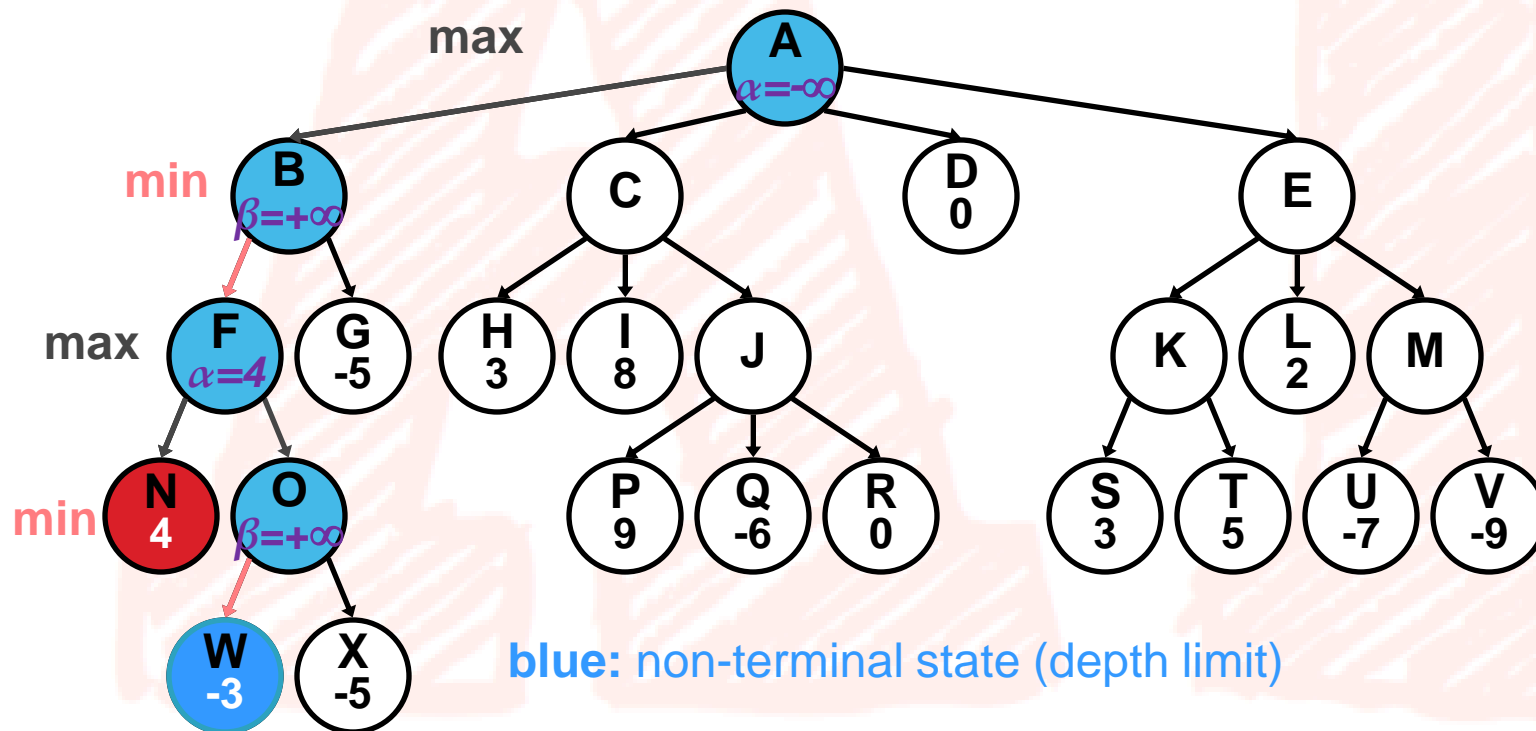
Call
Stack

O
F
B
A

Alpha-Beta Search Example

$\text{minimax}(W, 4, 4)$

evaluate and return SBE value



Call
Stack

W
O
F
B
A

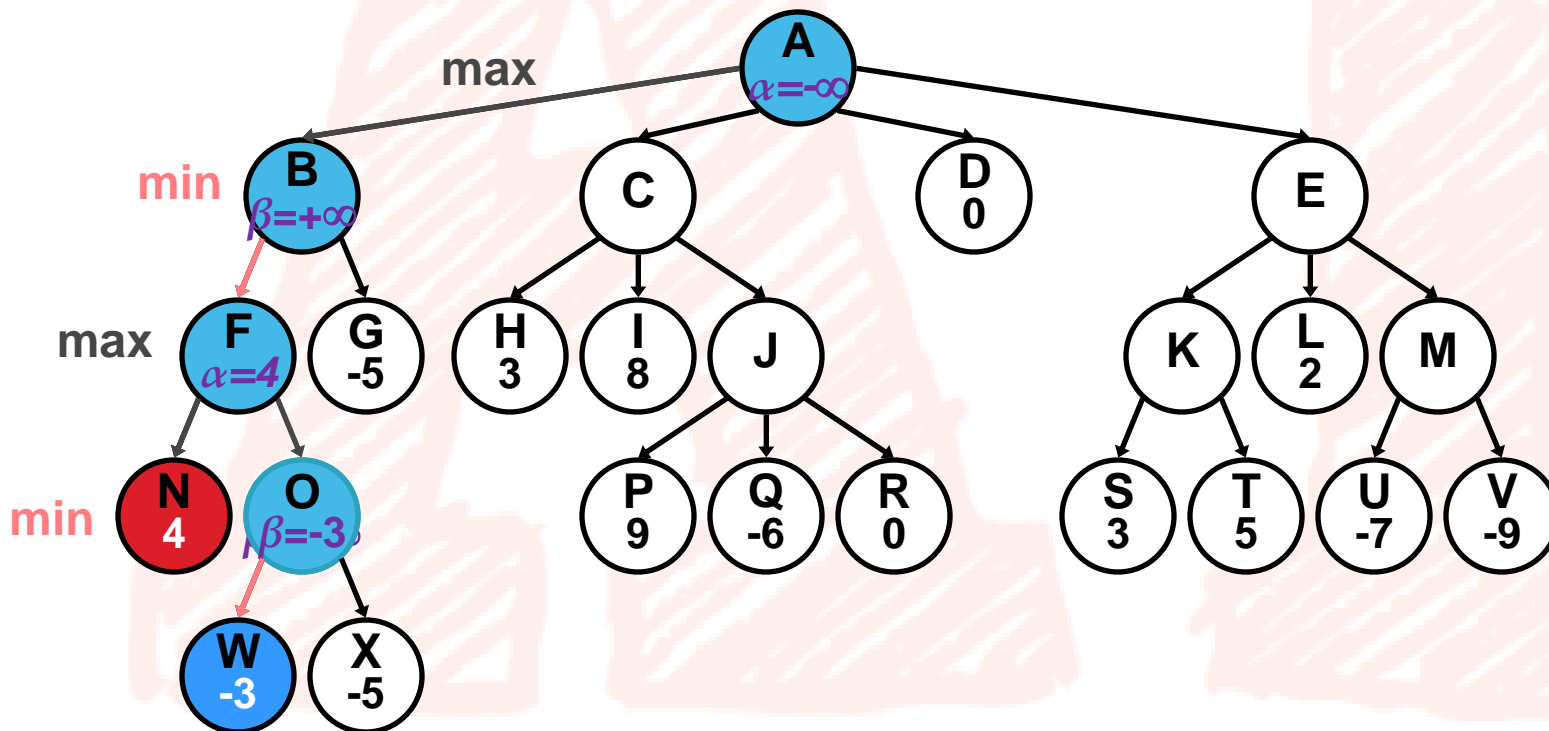
Alpha-Beta Search Example

back to
 $\text{minimax}(0, 3, 4)$

$\text{beta} = -3$, since $-3 \leq +\text{infinity}$ (minimizing)

Keep expanding O?

No since F's $\alpha \geq$ O's β is true: **alpha cutoff**



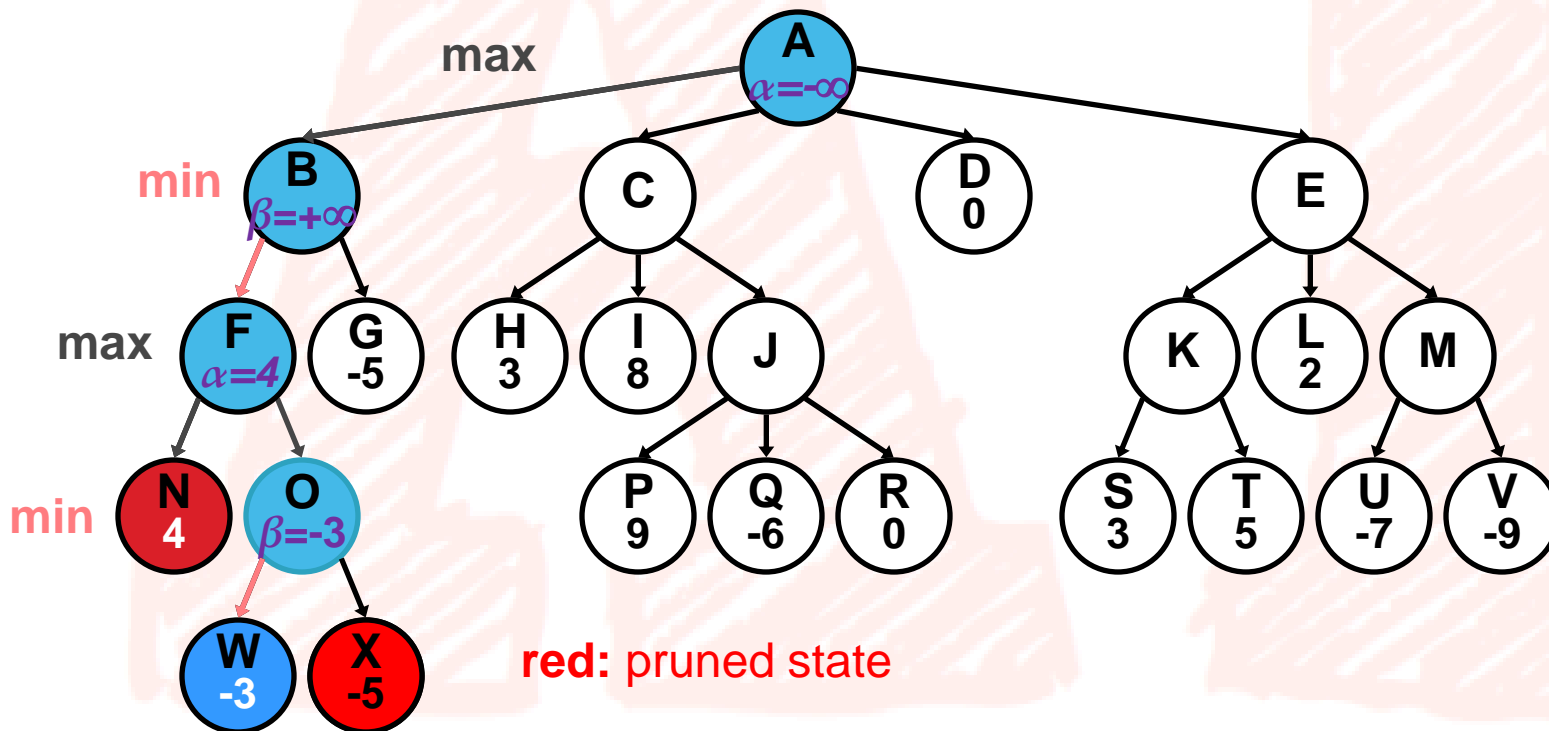
Call
Stack

O
F
B
A

Alpha-Beta Search Example

► Why?

- Smart opponent will choose W or worse, thus O's upper bound is -3.
Computer already has better move at N.

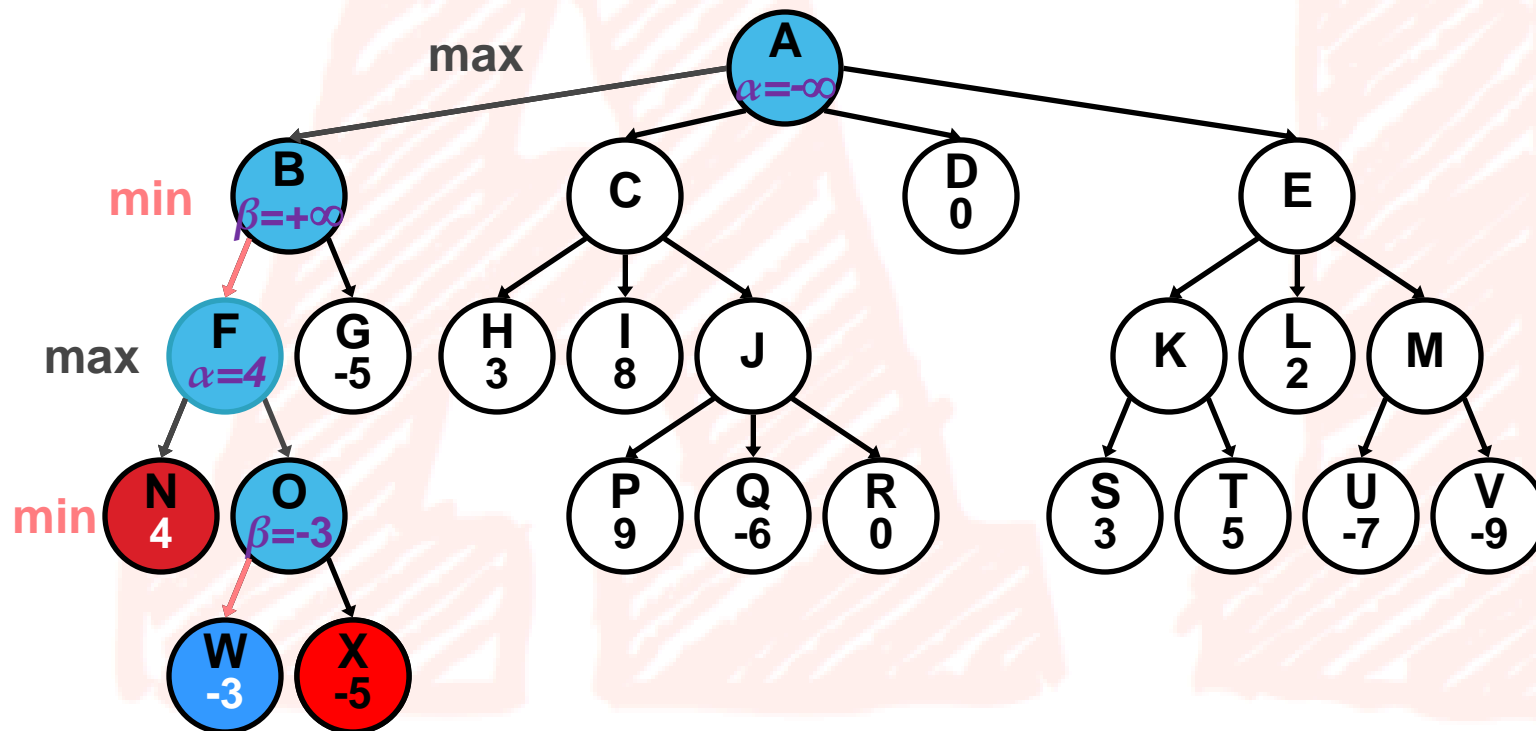


Alpha-Beta Search Example

back to
 $\text{minimax}(F, 2, 4)$

alpha doesn't change, since $-3 < 4$ (maximizing)

Keep expanding F? **No** since no more successors for F



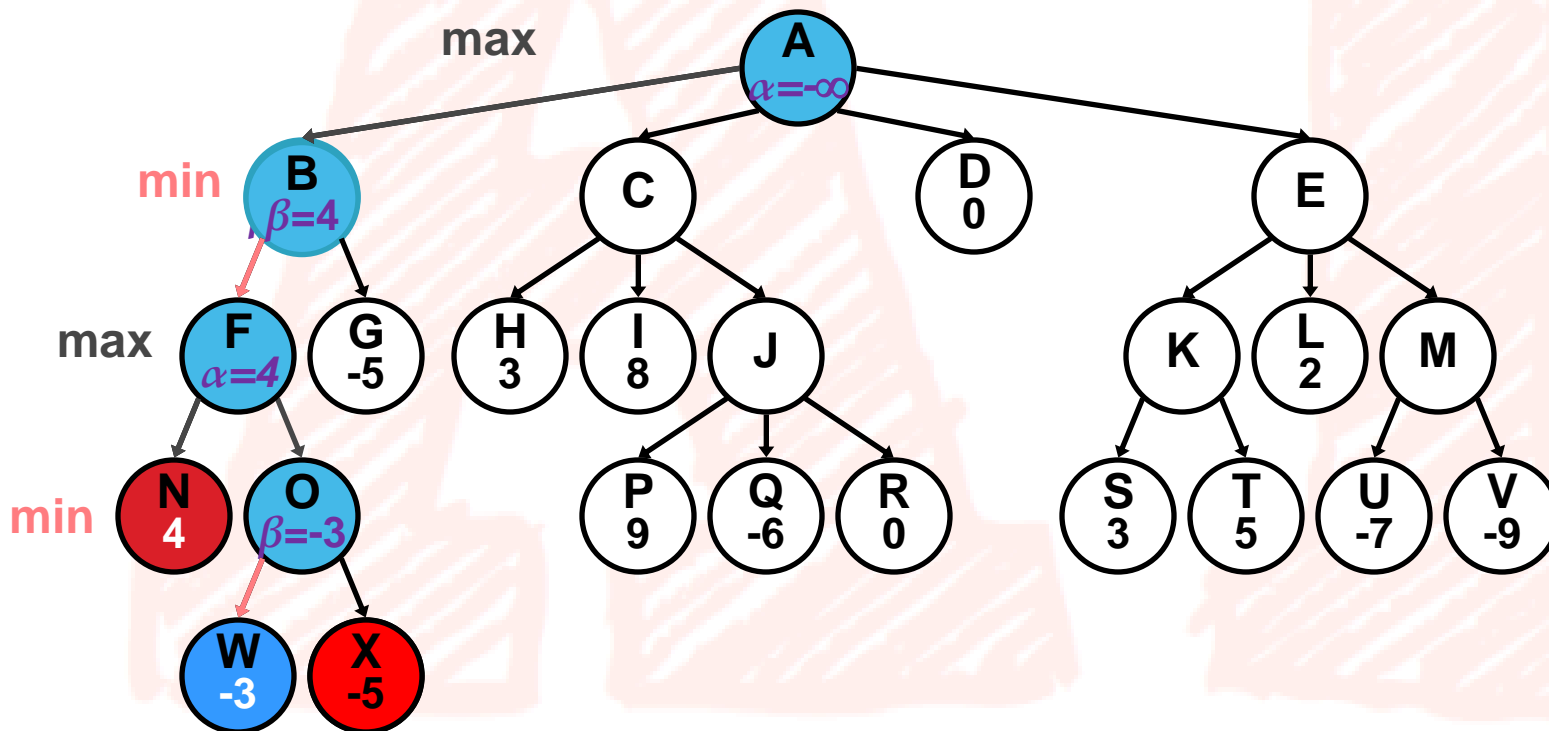
Call
Stack

F
B
A

Alpha-Beta Search Example

back to $\text{minimax}(B, 1, 4)$ $\text{beta} = 4$, since $4 \leq +\text{infinity}$ (minimizing)


Keep expanding B? Yes since A's $\alpha \geq B$'s β is **false**, no alpha cutoff



Alpha-Beta Search Example

$\text{minimax}(G, 2, 4)$

evaluate and return SBE value



Call
Stack

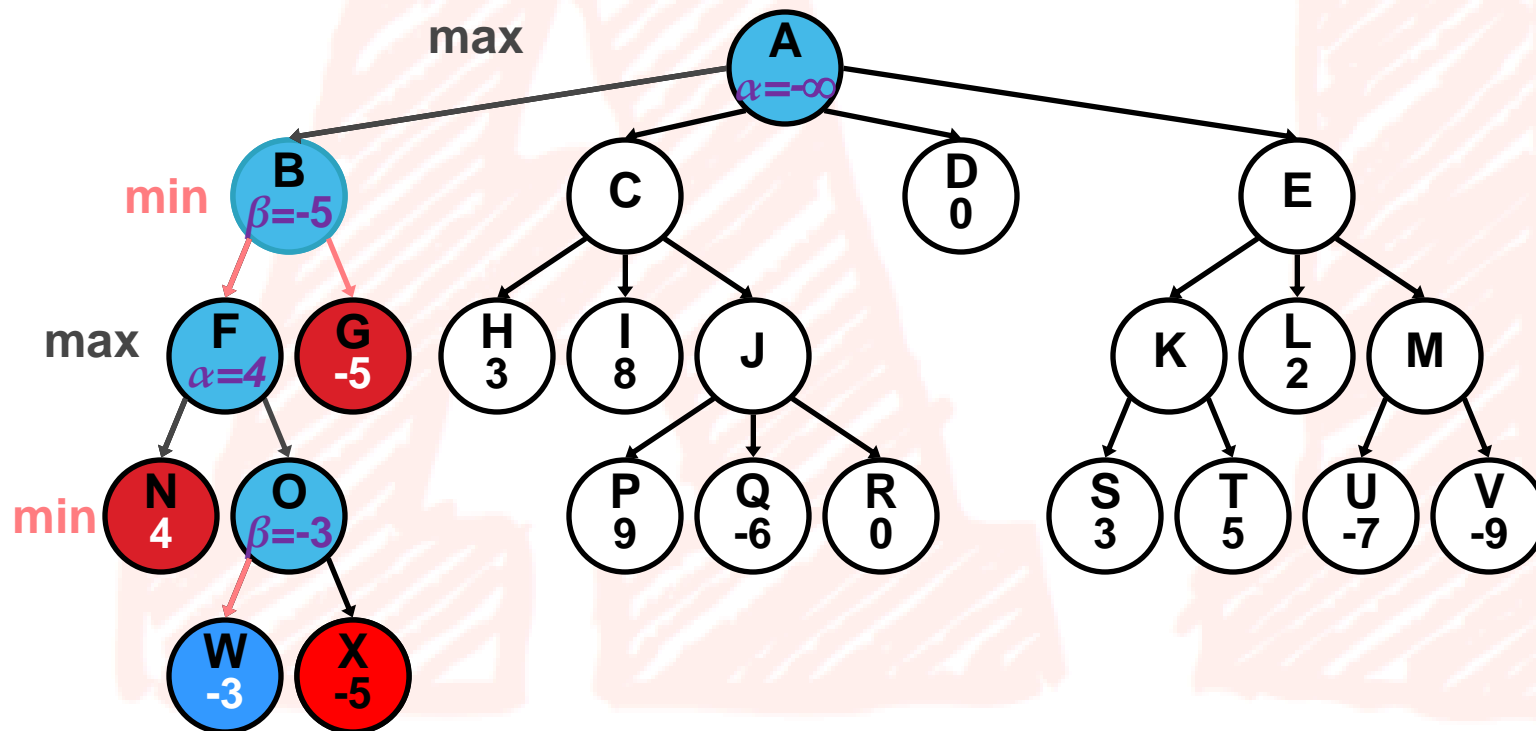
G
B
A

Alpha-Beta Search Example

back to
 $\text{minimax}(B, 1, 4)$

$\beta = -5$, since $-5 \leq 4$ (minimizing)

Keep expanding B? **No** since no more successors for B

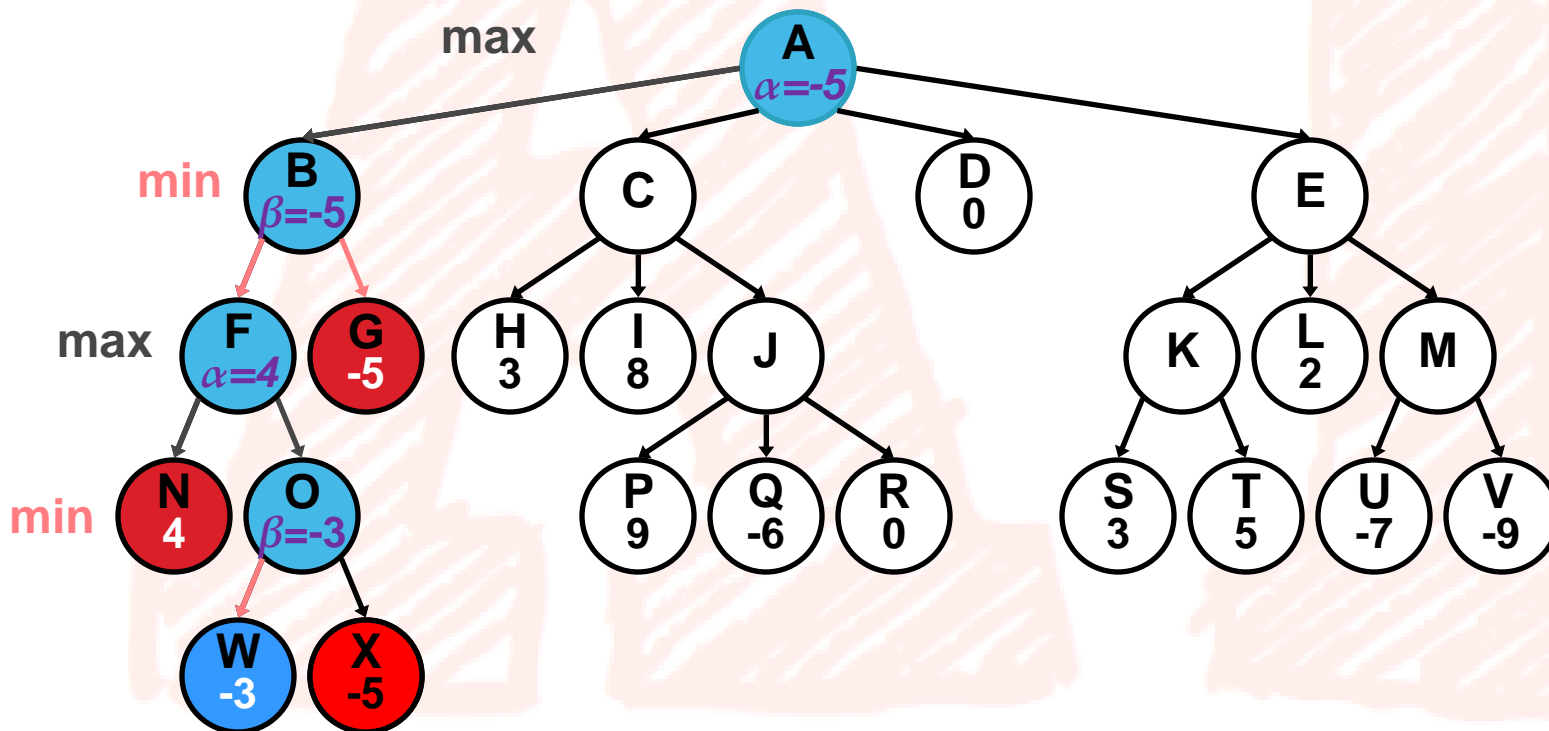


Alpha-Beta Search Example

back to
 $\text{minimax}(A, 0, 4)$

$\alpha = -5$, since $-5 \geq -\text{infinity}$ (maximizing)

Keep expanding A? Yes since there are more successors, no cutoff test



Call
Stack

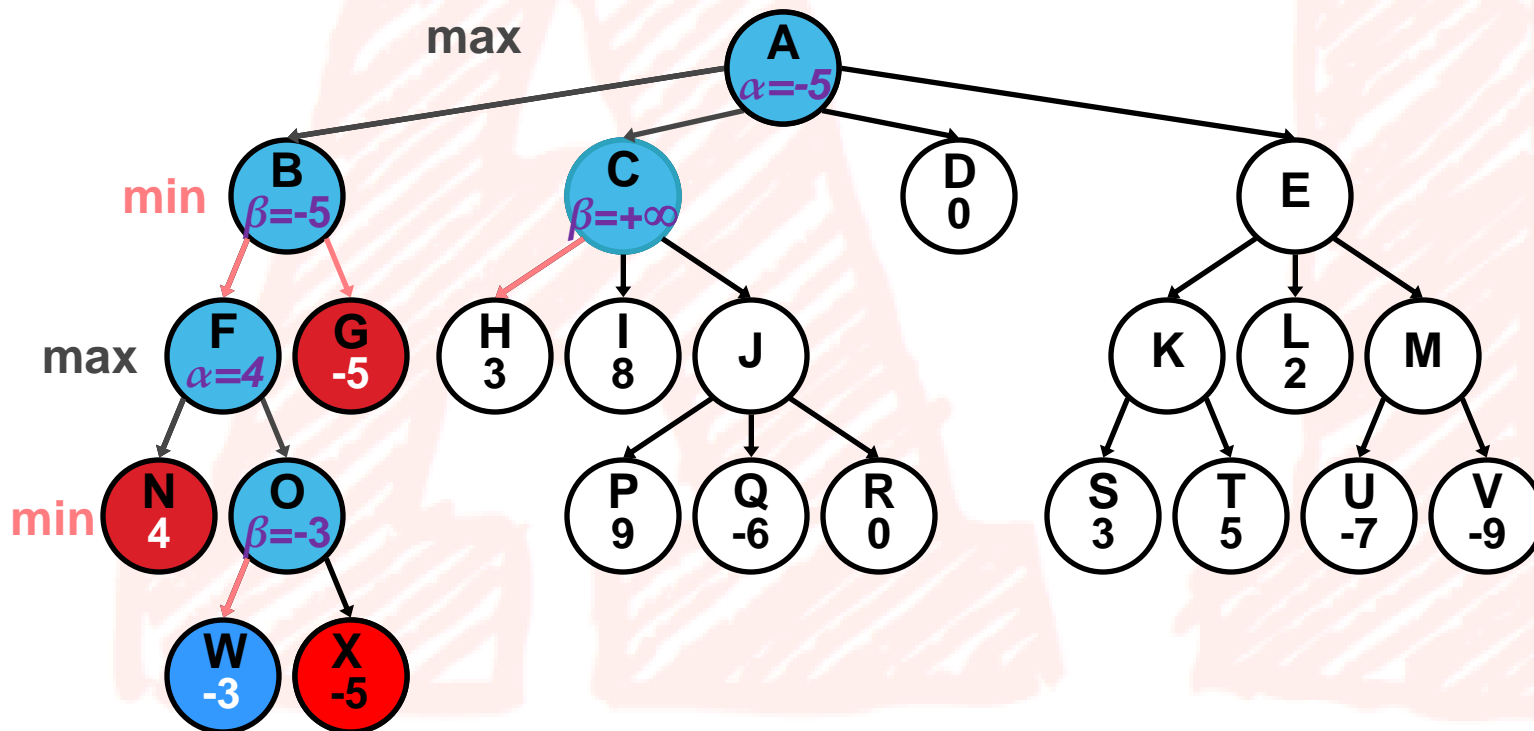
A

Alpha-Beta Search Example

$\text{minimax}(C, 1, 4)$

beta initialized to +infinity

Expand C? Yes since A's alpha \geq C's beta is **false**, no alpha cutoff



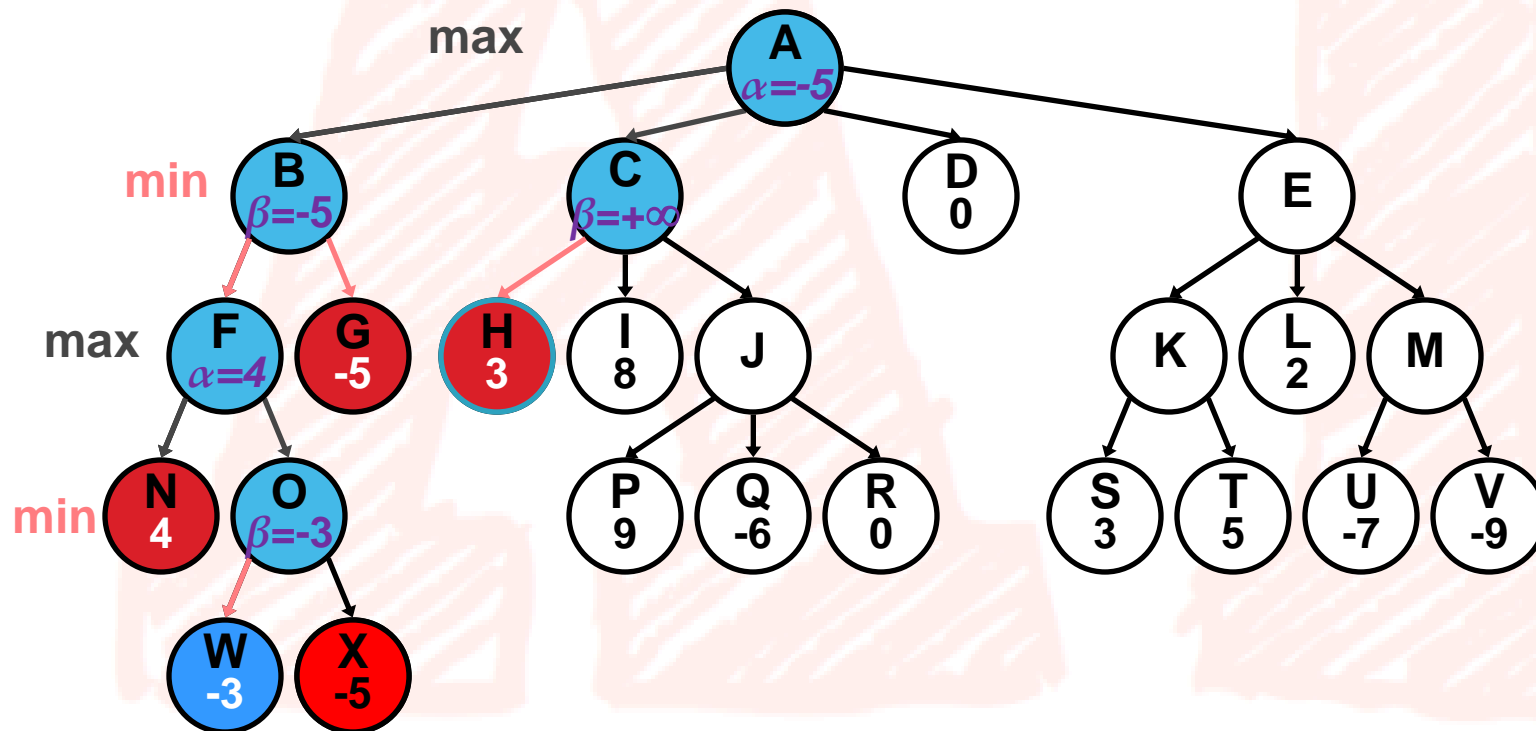
Call
Stack

C
A

Alpha-Beta Search Example

$\text{minimax}(H, 2, 4)$

evaluate and return SBE value



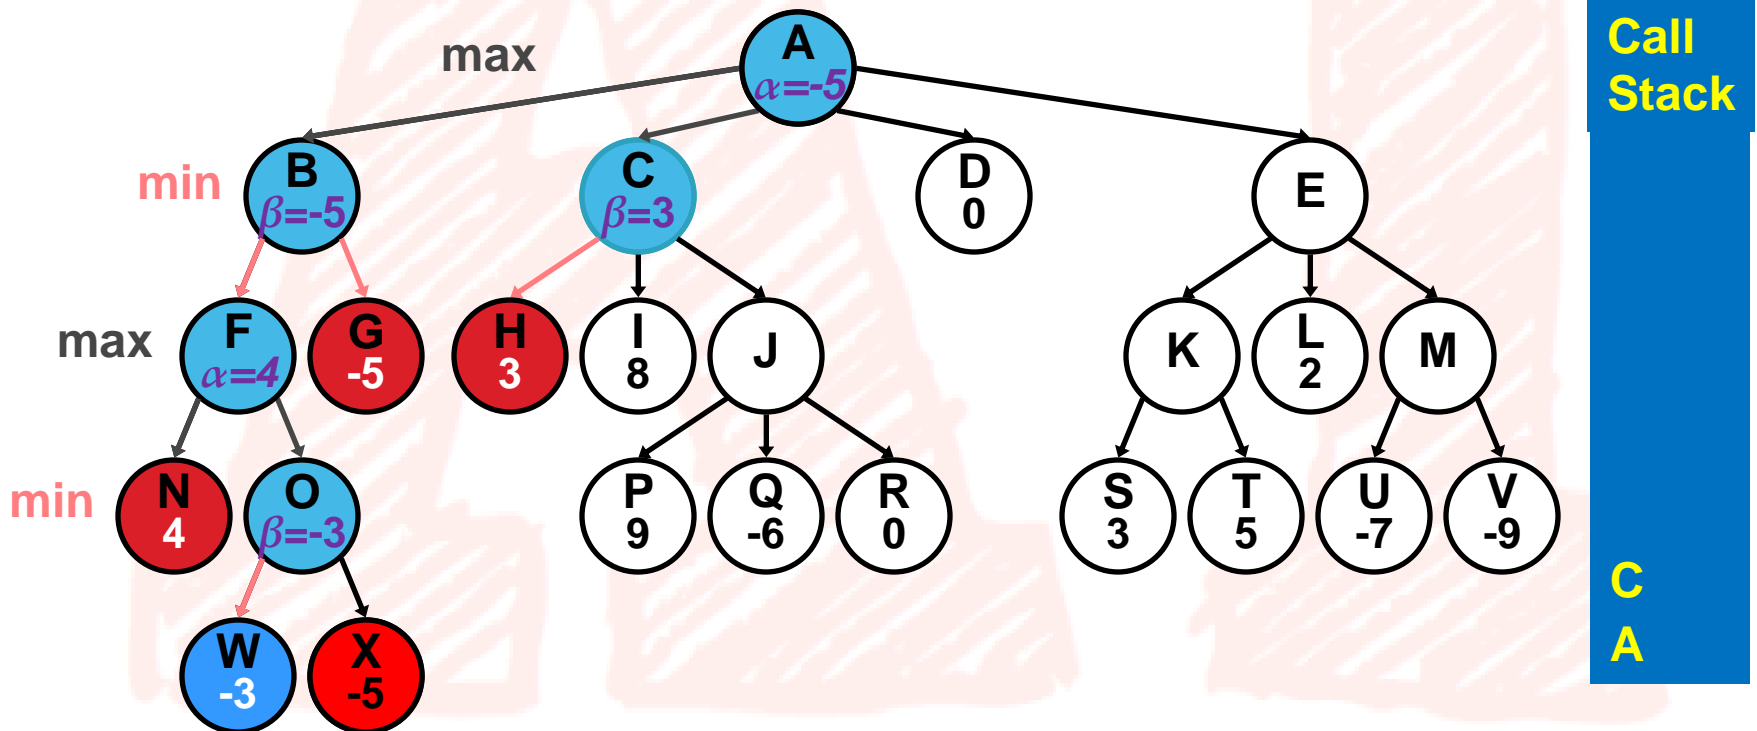
Call
Stack

H
C
A

Alpha-Beta Search Example

back to $\text{minimax}(C, 1, 4)$ $\text{beta} = 3$, since $3 \leq +\text{infinity}$ (minimizing)


Keep expanding C? Yes since A's $\alpha \geq C$'s β is **false**, no alpha cutoff



Alpha-Beta Search Example

minimax(I,2,4)

evaluate and return SBE value



Call Stack

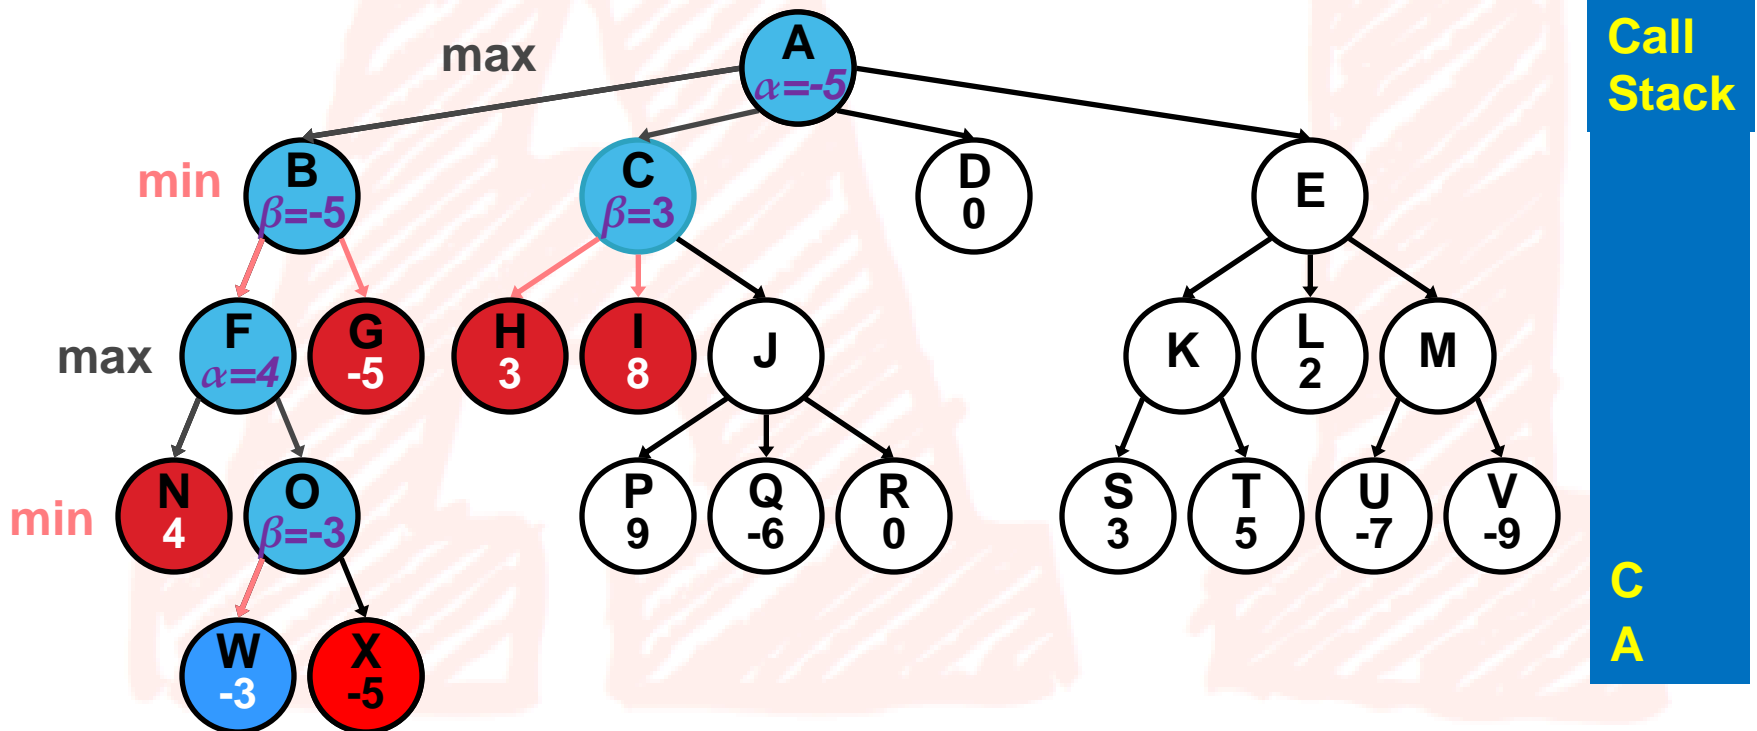
ICA

Alpha-Beta Search Example

back to
 $\text{minimax}(C, 1, 4)$

beta doesn't change, since $8 > 3$ (minimizing)

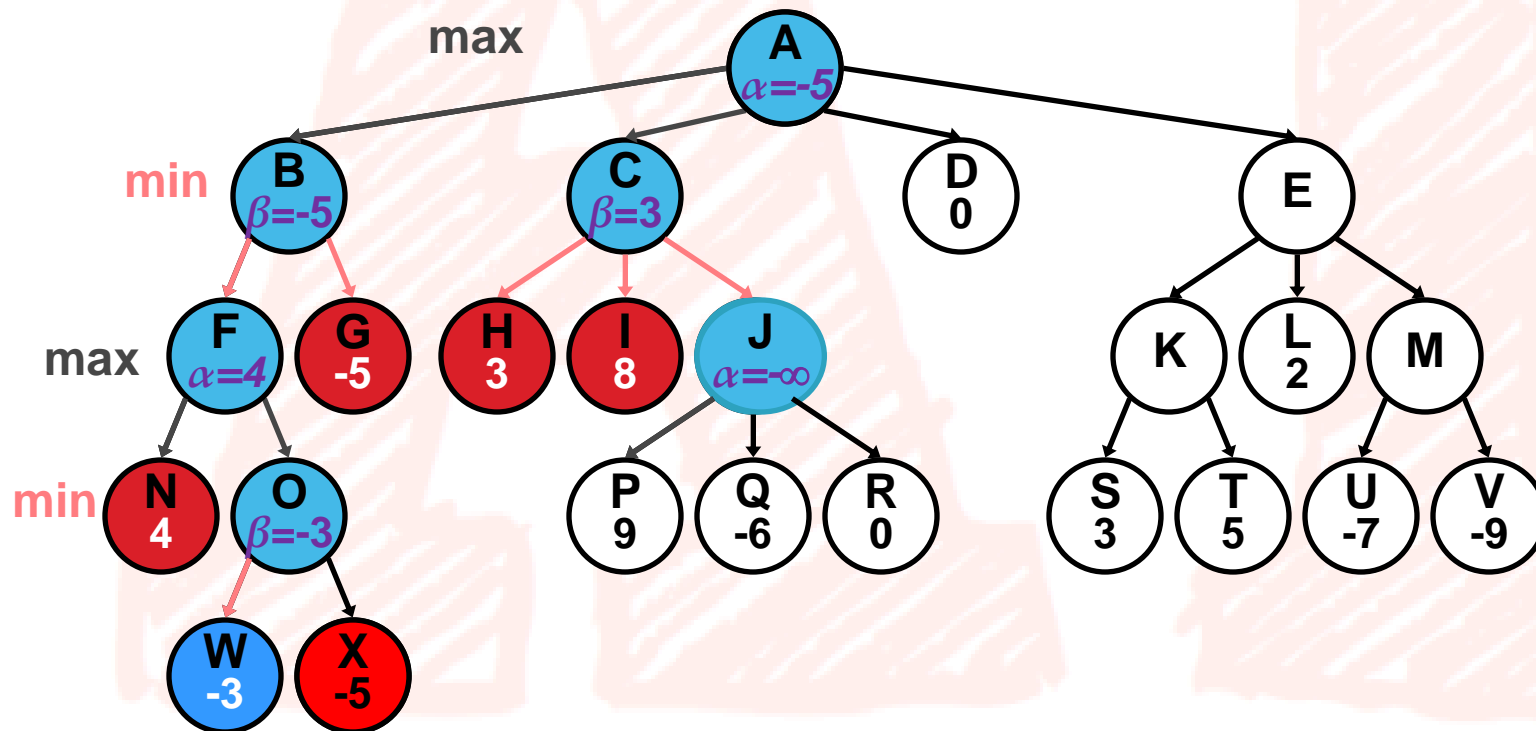
Keep expanding C? Yes since A's alpha \geq C's beta is **false**, no alpha cutoff



Alpha-Beta Search Example

$\text{minimax}(J, 2, 4)$ alpha initialized to -infinity

Expand J? Yes since J's alpha \geq C's beta is **false**, no beta cutoff



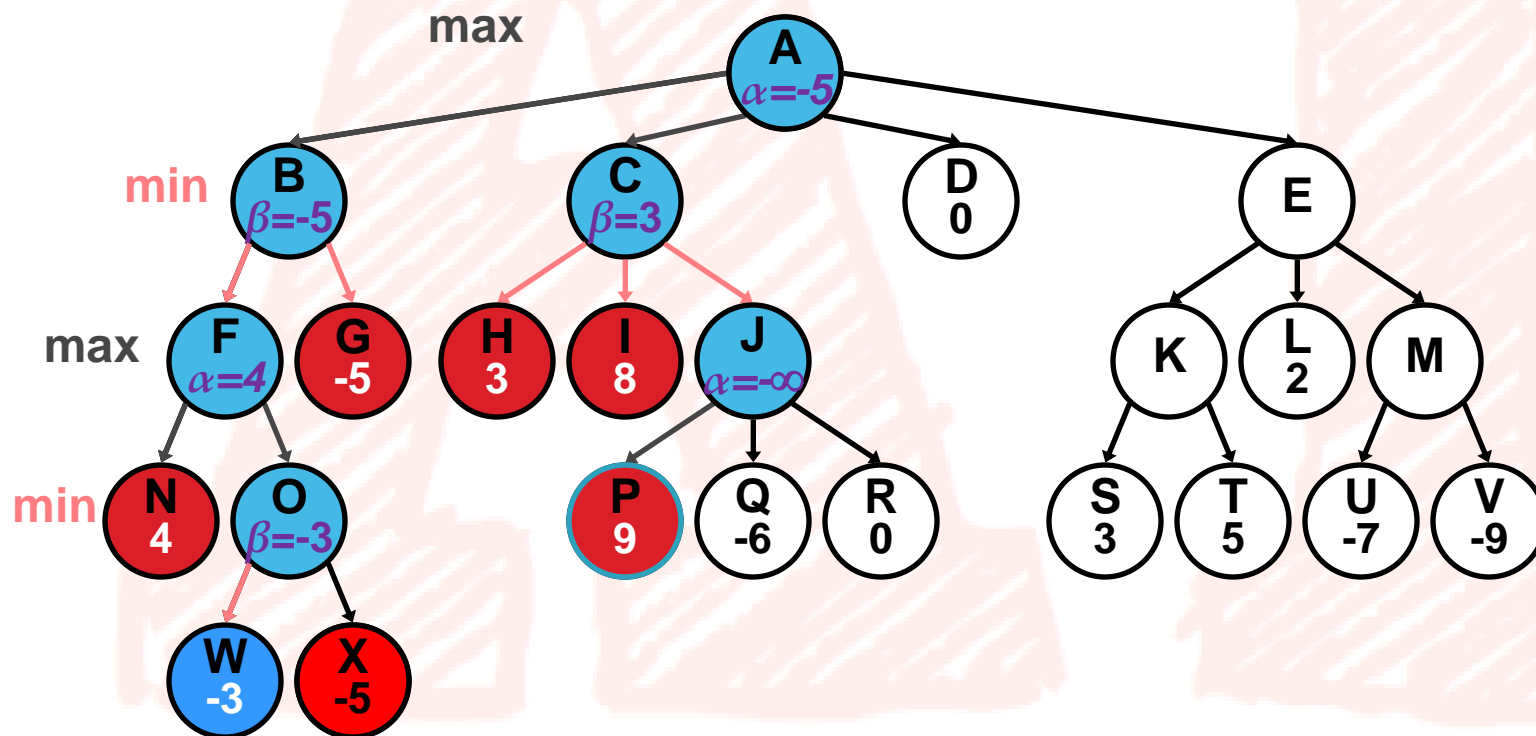
Call
Stack

J
C
A

Alpha-Beta Search Example

$\text{minimax}(P, 3, 4)$

evaluate and return SBE value



Call
Stack

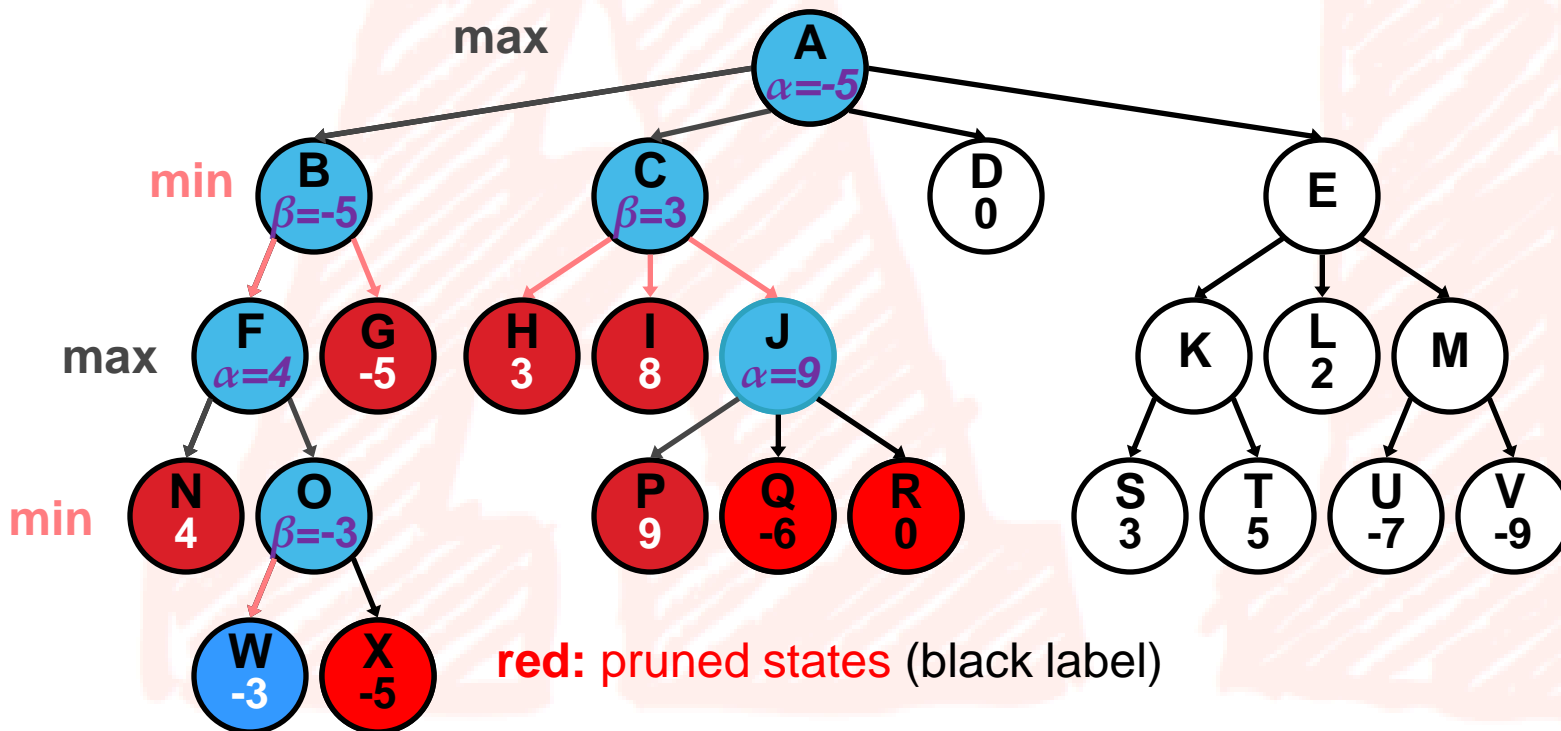
P
J
C
A

Alpha-Beta Search Example

back to
 $\text{minimax}(J, 2, 4)$

$\alpha = 9$, since $9 \geq -\text{infinity}$ (maximizing)

Keep expanding J? **No** since J's $\alpha \geq$ C's β is **true**: **beta cutoff**



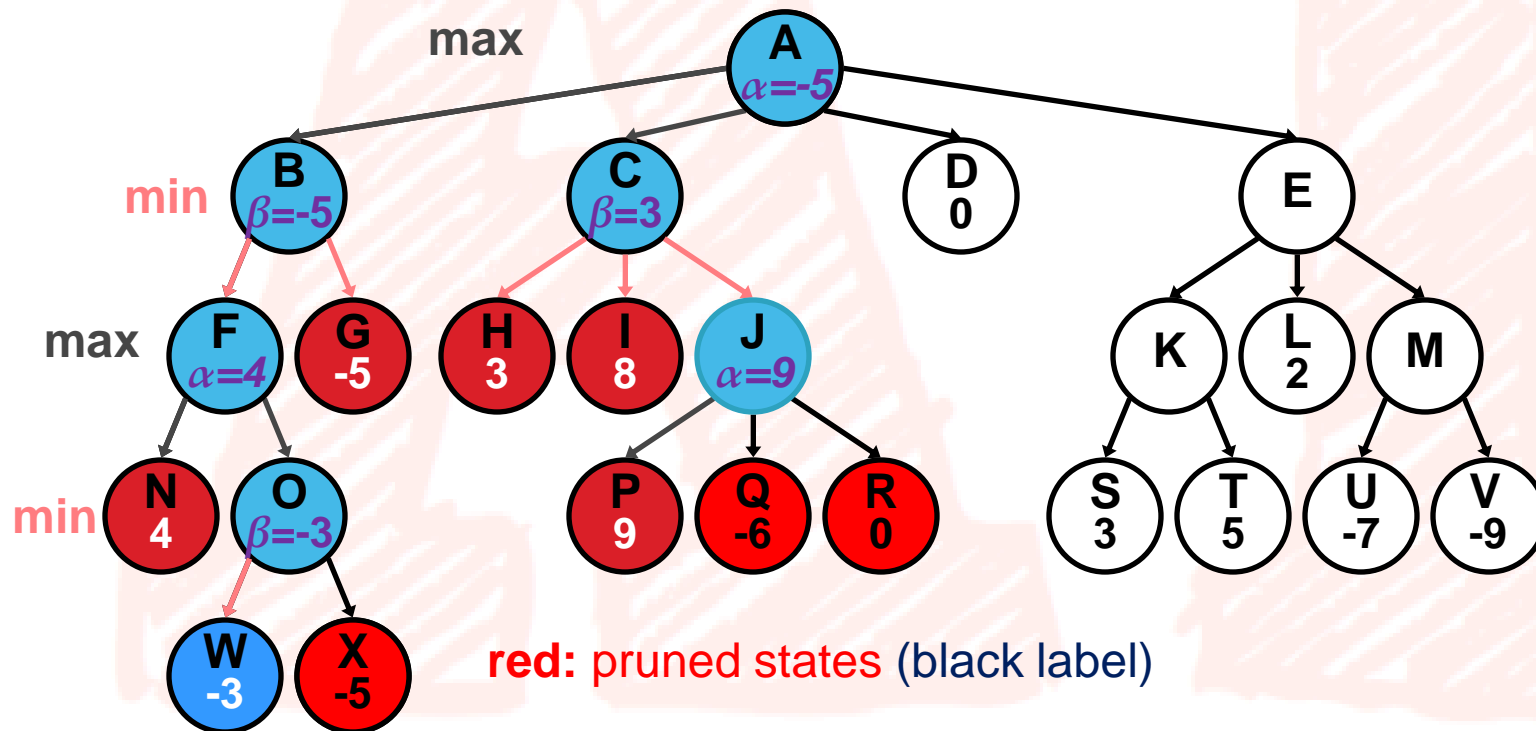
Call
Stack

J
C
A

Alpha-Beta Search Example

► Why?

- Computer will choose P or better, thus J's lower bound is 9.
Smart opponent won't let computer take move to J
(since opponent already has better move at H).



Call
Stack

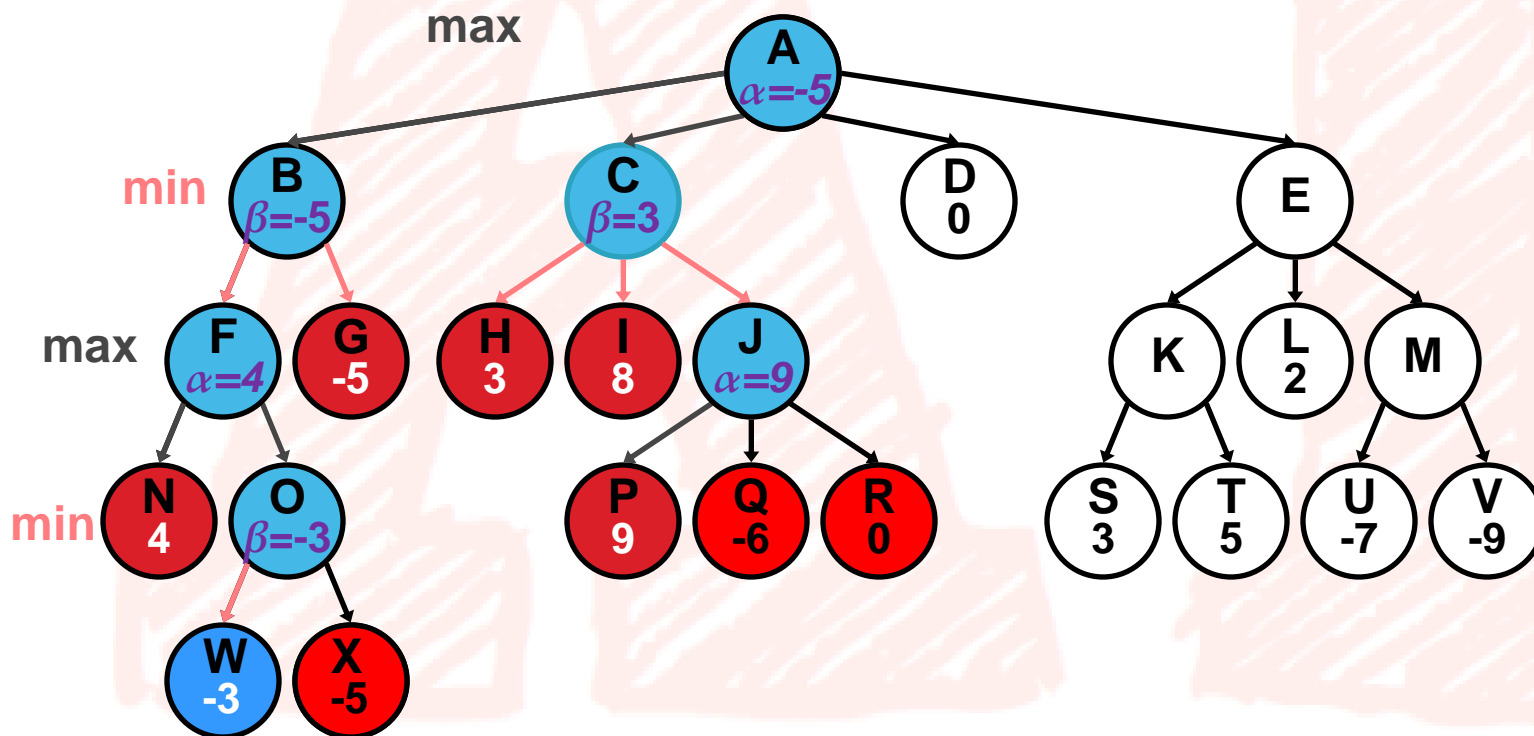
J
C
A

Alpha-Beta Search Example

back to
 $\text{minimax}(C, 1, 4)$

beta doesn't change, since $9 > 3$ (minimizing)

Keep expanding C? **No** since no more successors for C



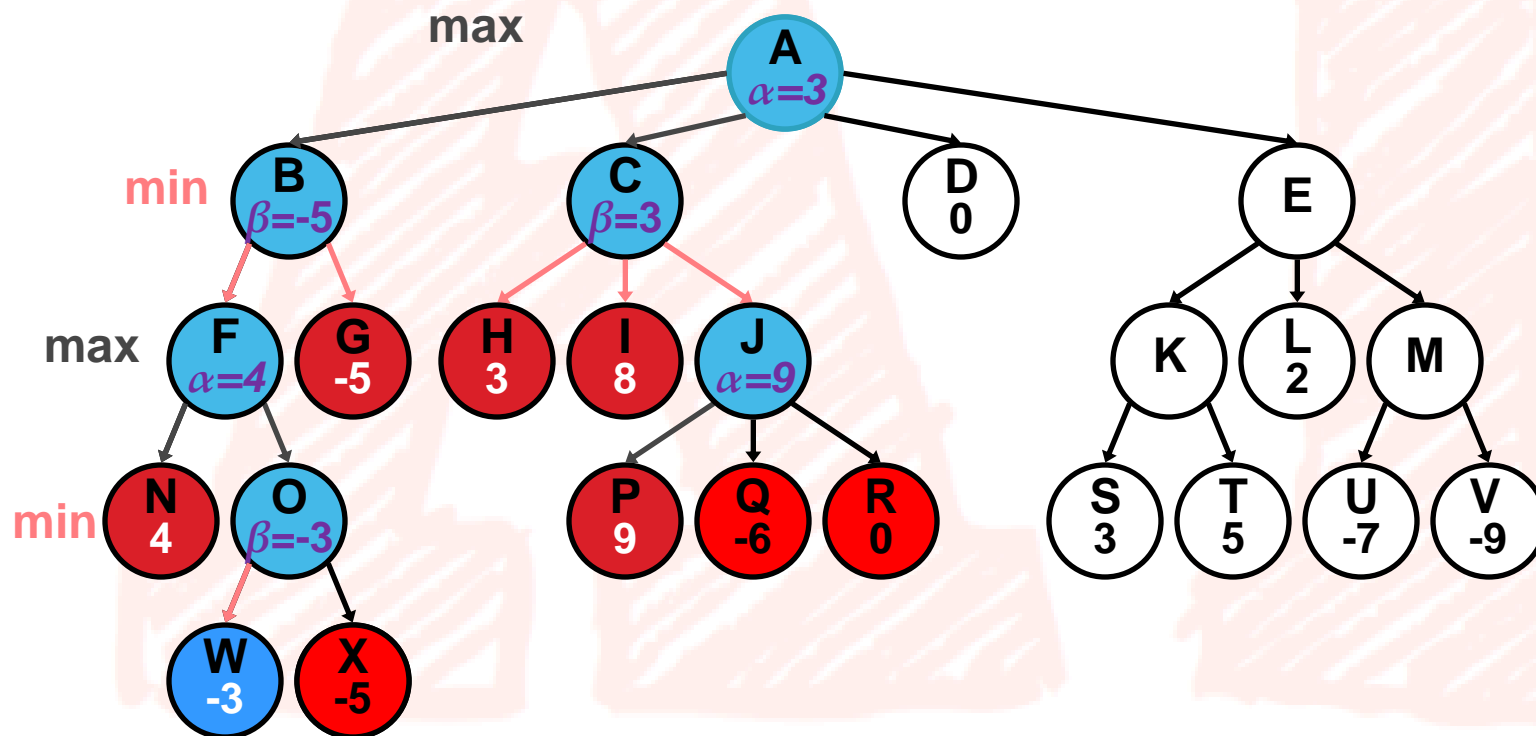
Call
Stack

C
A

Alpha-Beta Search Example

back to $\text{minimax}(A, 0, 4)$ $\alpha = 3$, since $3 \geq -5$ (maximizing)

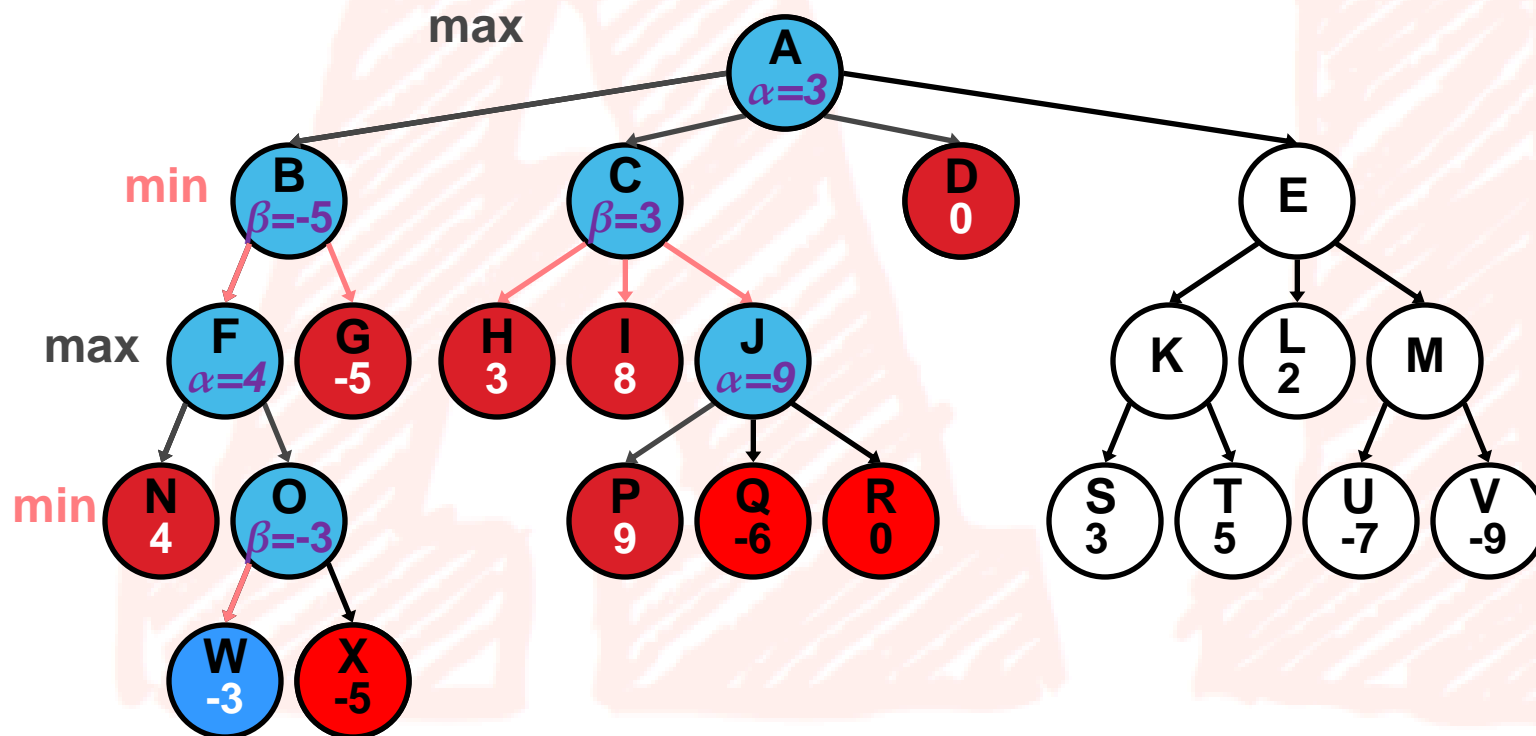
Keep expanding A? Yes since there are more successors, no cutoff test



Alpha-Beta Search Example

$\text{minimax}(D, 1, 4)$

evaluate and return SBE value



Call
Stack


D
A

Alpha-Beta Search Example

back to
 $\text{minimax}(A, 0, 4)$

alpha doesn't change, since $0 < 3$ (maximizing)

Keep expanding A? Yes since there are more successors, no cutoff test

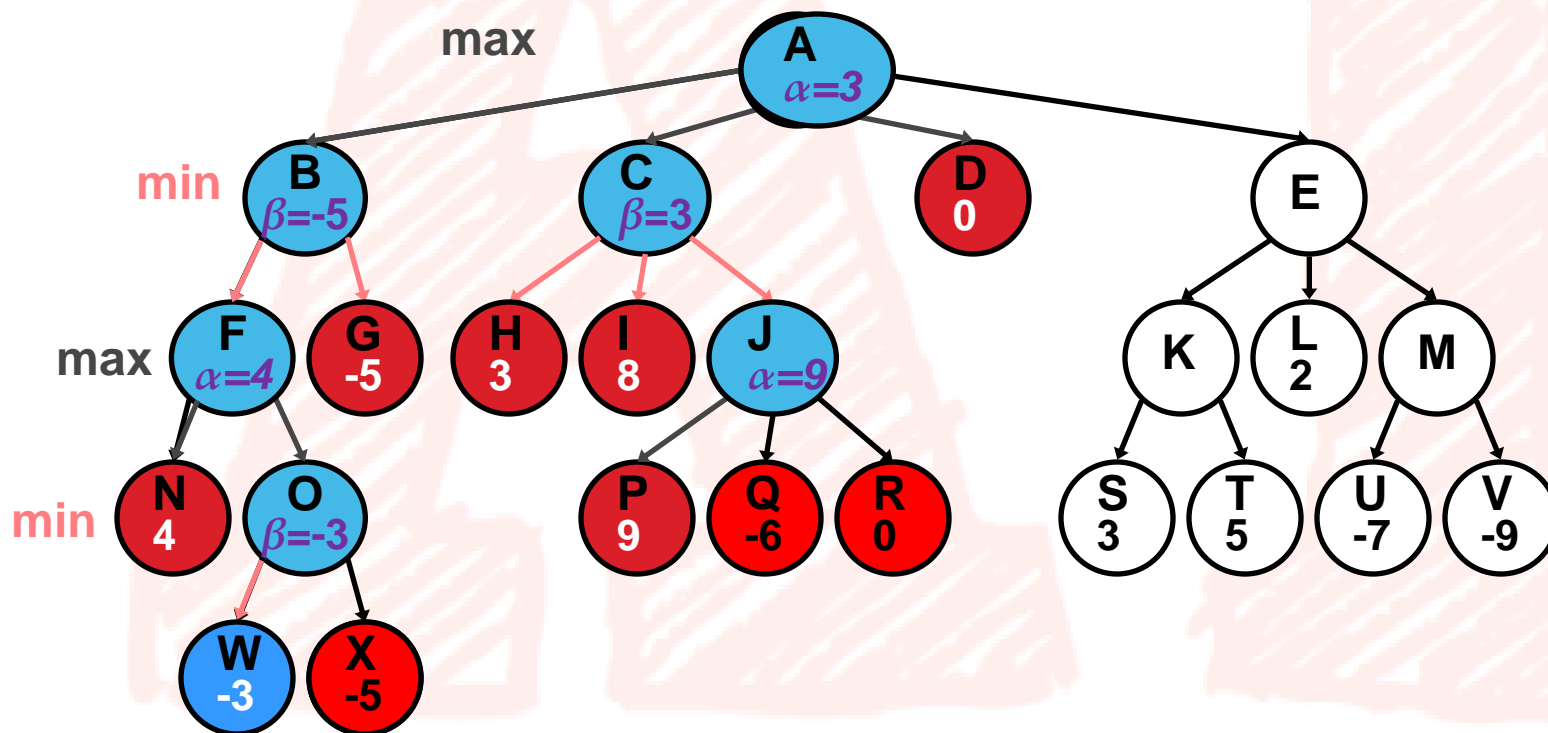


Call
Stack

A

Alpha-Beta Search Example

- How does the algorithm finish searching the tree?




Call
Stack

A

Alpha-Beta Search Example

- ▶ Stop Expanding E since A's alpha \geq E's beta is true: alpha cutoff
- ▶ Why?
 - Smart opponent will choose L or worse, thus E's upper bound is 2.
 - Computer already has better move at C.




Call
Stack

A

Alpha-Beta Search Example

- Result: Computer chooses move to C.




Call
Stack

A


Exercise 3

- Perform the minimax algorithm on the figure below with **Alpha-Beta-pruning**




Exercise 4

- ▶ Perform the minimax algorithm on the figure below with Alpha-Beta-pruning



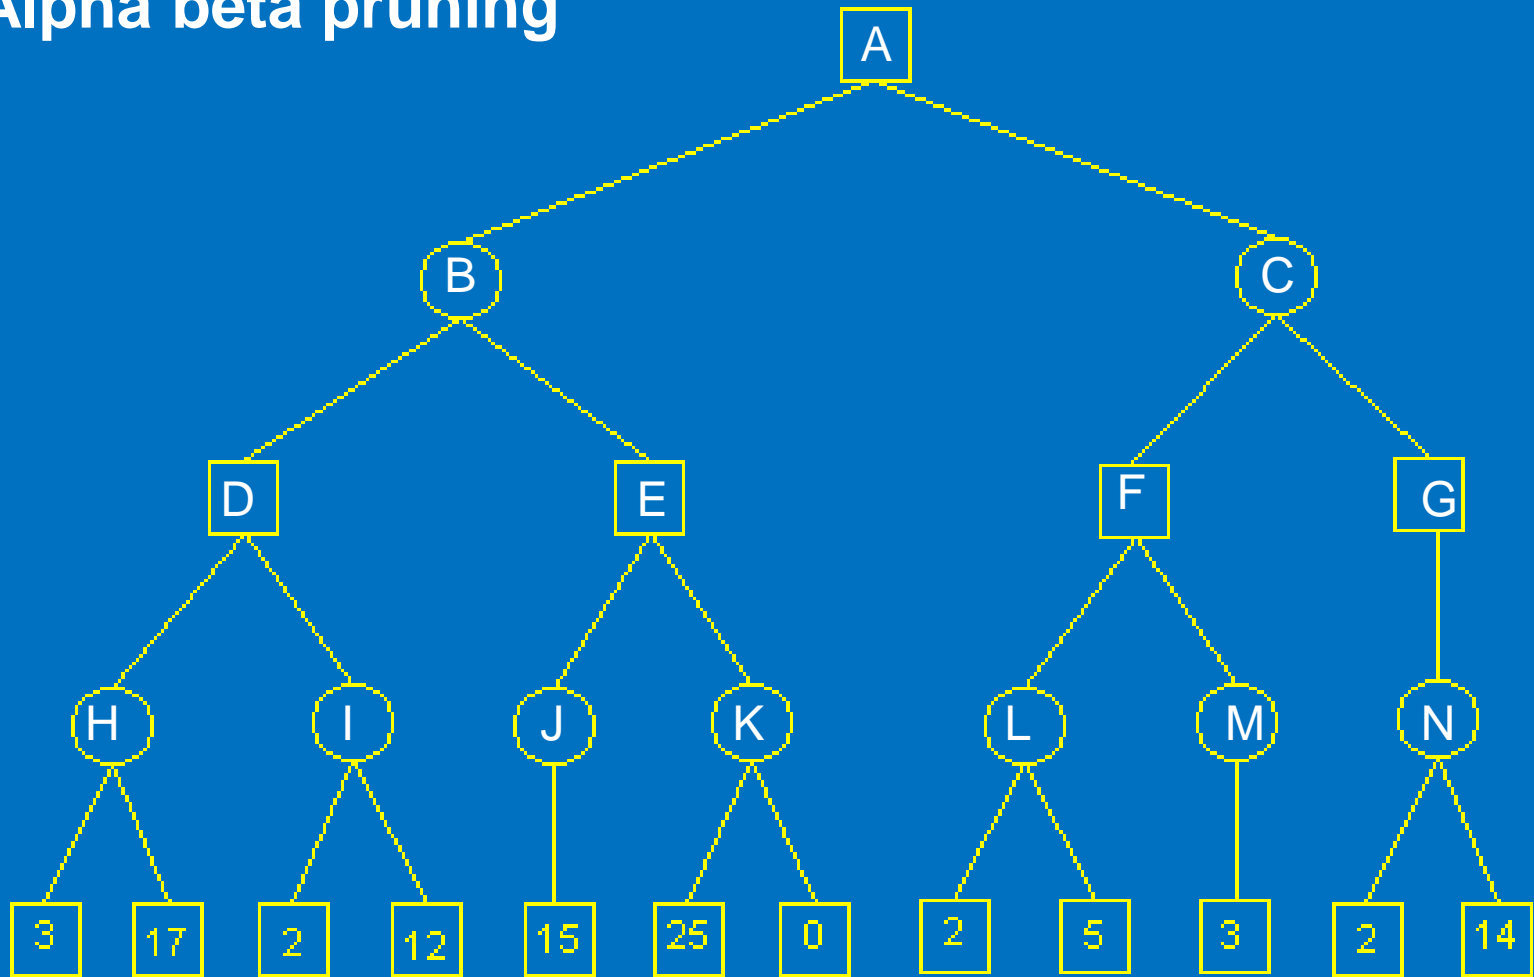
Exercise 5

- Perform the minimax algorithm on the figure below with Alpha-Beta-pruning




Exercise 6

Alpha beta pruning




Exercise 6: Result

Alpha beta pruning



Exercise 7



Alpha-Beta Algorithm

function ALPHA-BETA-SEARCH(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value v

function MAX-VALUE(*state*, α , β) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for each s **in** SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$


return v

Alpha-Beta Algorithm

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for each  $s$  in SUCCESSORS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```

General alpha-beta pruning

- ▶ Consider a **node n** somewhere in the tree
- ▶ If player has a **better choice at**
 - Parent node of **n**
 - Or any choice point further up
- ▶ **n** will never be reached in actual play.
- ▶ Hence when enough is known about **n** , it can be pruned.



Final Comments: Alpha–Beta Pruning

- ▶ Pruning does not affect final results
- ▶ Entire subtrees can be pruned.
- ▶ Good *move ordering* improves effectiveness of pruning
- ▶ With "perfect ordering" time complexity is $O(b^{m/2})$
 - Branching factor of \sqrt{b} !!
 - Alpha–beta pruning can look twice as far as Minimax in the same amount of time
- ▶ Repeated states are again possible.
 - Store them in memory = transposition table

Imperfect Real-Time Decisions

- ▶ Minimax require **too much leaf-node evaluations**.
- ▶ May be impractical within a reasonable amount of time.
- ▶ SHANNON (1950):
 - Cut off search earlier (**replace TERMINAL-TEST by CUTOFF-TEST**)
 - Apply heuristic evaluation function EVAL (replacing utility function of alpha-beta)

Cutting off search

- ▶ Change:

```
if TERMINAL-TEST(state) then  
  return UTILITY(state)
```

into




```
if CUTOFF-TEST(state, depth) then  
  return EVAL(state)
```

- ▶ Introduces a fixed-depth limit depth
 - Is selected so that the amount of time will not exceed what the rules of the game allow.
- ▶ When cutoff occurs, the evaluation is performed.

Heuristic EVAL

- ▶ EVAL function returns an estimate of **the expected utility of the game from a given position**.
- ▶ Performance of game playing depends on quality of **EVAL**.
- ▶ Requirements:
 - **EVAL must agree with terminal-nodes** in the same way as UTILITY.
 - Computation **may not take too long**.
 - For non-terminal states the EVAL should be strongly correlated with the actual chance of winning.
- ▶ Only useful for quiescent (no wild swings in value in near future) states


Heuristic **EVAL** example




- ▶ Heuristic: $E(n) = M(n) - O(n)$
 - $M(n)$: total win paths of X,
 - $O(n)$: total win paths of O

Heuristic EVAL example


$$Eval(s) = w_1f_1(s) + w_2f_2(s) + \dots + w_nf_n(s)$$




(a) White to move
Fairly even



(b) Black to move
White slightly better



(c) White to move
Black winning



(d) Black to move
White about to lose

- w_i : a weight,
- f_i : a feature of the position

- Pawn: 1,
- Knight: 3,
- Rook: 5,
- Queen: 9

Multiplayer games

- ▶ Games allow **more than two players**
- ▶ Single **minimax values become vectors**

to move

A

(1, 2, 6)

B

(1, 2, 6)

(-1, 5, 2)

C

(1, 2, 6)

X

(6, 1, 2)

(-1, 5, 2)

(5, 4, 5)

A

(1, 2, 6)

(4, 2, 3)

(6, 1, 2)

(7, 4, -1)

(5, -1, -1)


(-1, 5, 2)

(7, 7, -1)

(5, 4, 5)

Games that include chance

Backgammon



MAX


CHANCE

MIN

CHANCE

MAX



TERMINAL



- Possible moves (5-10, 5-11), (5-11, 19-24), (5-10, 10-16) and (5-11, 11-16)


Games that include chance

Backgammon



- ▶ Possible moves (5-10, 5-11), (5-11, 19-24), (5-10, 10-16) and (5-11, 11-16)
- ▶ [1,1], ..., [6,6] chance $1/36$, all other chance $1/18$

Backgammon



- expected value

Expected minimax value

EXPECTED-MINIMAX-VALUE(n) =

UTILITY(n)

If n is a terminal


$\max_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$ If n is a max node

$\min_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$ If n is a min node

$\sum_{s \in \text{successors}(n)} P(s) \cdot \text{EXPECTEDMINIMAX}(s)$ If n is a
chance node

- ▶ These equations can be backed-up recursively all the way to the root of the game tree.

Position evaluation with chance nodes



- ▶ Left, A₁ wins
- ▶ Right, A₂ wins
- ▶ **Outcome of evaluation function may not change** when values are scaled differently.
- ▶ Behavior is preserved only by a positive linear transformation of EVAL.

Summary

- ▶ Games are fun (and dangerous)
- ▶ They illustrate several important points about AI
 - Perfection is unattainable \rightarrow approximation
 - Good idea what to think about
 - Uncertainty constrains the assignment of values to states
- ▶ Games are to AI as grand prix racing is to automobile design.



FACULTY OF INFORMATION TECHNOLOGY

Thank you for your attention!